

SPECIFYING AND CHECKING CORRECTNESS PROPERTIES

with Linear Temporal Logic (LTL)

safety and liveness

Safety properties...

... express behavior that a system should avoid

Something bad cannot happen!

Modeler's
perspective

Safety properties... aka state properties

... express behavior that a system should avoid

Safety properties are about **states**...

- **absence of deadlocks**
- **system invariants** that hold
 - *globally*: in all system states
 - *locally*: in selected local states, always



Verifier's
(Model Checker's)
perspective

Liveness properties...

... express behavior that the system should allow

Something good should happen!

- progress, termination, non-starvation, fairness,
or (usually) some other kind of positive
eventuality

Modeler's
perspective

Liveness properties... aka path properties

... *express behavior that the system must allow

Liveness properties are about **paths**...

... combining state properties along an execution path to express dependencies between them

Verifier's
(Model Checker's)
perspective

Specifying Properties in Promela/Spin



State properties

- assertions
- end-state labels
- Linear Temporal Logic (LTL) properties (only with [])

Path properties

- Linear Temporal Logic (LTL) properties (containing <>, U, X)
- never claims
- progress-state labels (desirable cycles)

Assertions are...

... the simplest way to express safety properties

`assert(boolean_expression)`
• State proposition

- always executable
- can be used anywhere a statement can

Assertion can define a local invariant

When used inside any process

```
proctype receiver() {  
    ...  
    toReceiver ? msg;  
    assert(msg != ERROR); // must be true when this process reaches this state  
    ...  
}
```

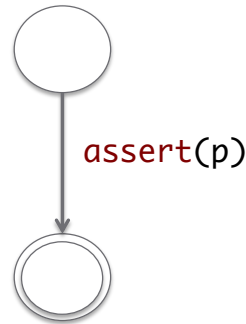
Violation anytime during verification will produce an error with a corresponding failure trail

A local invariant is a property that always holds true in specific local states

Assert can be used to check a global invariant

Add a monitor, or *spy*, process to the model, and run

```
active proctype monitor()  
{  
    assert(p) ;  
}
```



`assert(p)` will be enabled at every global state, and checked by verifier

A **global invariant** is a property that holds true in every global state (global state = a state of the product automaton)

Deadlock: being stuck in an invalid end state

- Automatically detected by default in Spin in Safety verification
 - to disable, don't check *Invalid end states* in SpinRCP
- The closing curly brace of a process is a **valid** end state

End-state labels define legal termination

To indicate other valid end-states, use state labels prefixed with **end**:

end: ...

end1: ...

endstate: ...

end-state: ...



A system that stops : Sema1.pml

```

mtype { P, V }

chan s = [0] of { mtype }

active proctype Sema() {
    bit free = 1;

    do
        :: (free == 1) -> s!P; free = 0;
        :: (free == 0) -> s?V; free = 1;
    od;
}

active [3] proctype User() {
    s?P; // enter, no reservation
    critical: skip;
    s!V; // leave
}
  
```

rendez-vous channel ensures
atomicity of semaphore
operations,
which must be indivisible
(atomicity guaranteed when
implemented by an OS)



- Run *Safety Verification*
 - *Invalid end states*
- *Check that it fails due to invalid end state*
- *But it should be valid...*
- Where do you put the end state in Sema?
 - Before (free == 1)?
 - Before s?V
 - Before s!P



Can fix safety check by
adding an end state

Sema2.pml

Source: public-spin/Semaphore

```
mtype { P, V }  
  
chan s = [0] of { mtype }  
  
active proctype Sema() {  
    bit free = 1;  
  
    do  
        :: (free == 1) ->  
            end: s!P; free = 0;  
        :: (free == 0) ->  
            s?V; free = 1;  
    od;  
}  
  
active [3] proctype user() {  
    s?P;  
    critical: skip;  
    s!V;  
}
```

end state can't be here because
one of the do guards is always
executable

this is a valid end state of
the semaphore: waiting for
another P instruction

- Dark the releasing of the lock as a valid end state. It will pass the verification.

Now each user process enters its
critical section only once and
terminates: an end state is needed in
semaphore implementation to tell
Spin this is ok during a safety check



Sema3.pml - Try at home!

```
mtype { P, V }  
  
chan s = [1] of { mtype };  
  
active proctype Dijkstra() {  
    bit free = 1;  
  
    do  
        :: (free == 1) ->  
            end: s!P; free = 0;  
        :: (free == 0) ->  
            s?V; free = 1;  
    od;  
}  
  
active [3] proctype user() {  
    s?P;  
    critical: skip;  
    s!V;  
}
```

- Replace rendez-vous channel with a regular channel • *It's not atomic now. Will not pass.*
- Does the verification still pass?
- If not, why?
- Inspect content of chan s: what does it contain?

See Appendix for more examples on the effect of rendez-vous!

Linear Temporal Logic

Can express global correctness properties explicitly and compactly

– both **state** and **path**

- *Logic*: a formal logic in the full mathematical sense
- *Temporal*: operates on ordered execution sequences
- *Linear*: time flows along a straight line, no branching

Linear Temporal Logic formulas are implicitly
universally quantified over execution paths

\forall

“for all paths (reachable execution traces) of the system”

vs.

~~“there exists a path in the system”~~

An LTL formula consists of temporal operators...

p: a state proposition (*in Promela: boolean expression that can be checked on a system state, like something we can put in an assert(...) statement*)
– means p is true at the current state

Always $[]p$ p is true throughout

Eventually $<>p$ p is eventually true

Next Xp p is true in the next state

Until (Strong) $p_1 \cup p_2$ p_1 stays true until p_2 becomes true

Main
temporal
operators

... and logical operators: $||$, $\&\&$, $!$, $->$, $<->$

• Means implication

• Negation

$->$ has a different meaning in LTL;
it's implication, not a statement
separator

What is a state proposition?



Any valid PROMELA Boolean expression:

$(\text{mutex} < 2)$

$(x > 0 \ \&\& \ y < 0)$

$(\text{len}(\text{ch}) > 0)$

$(b == 0)$



...

LTl properties are easy to visualize

*for a
single
path*

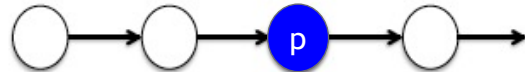
p : p holds in current state



Xp : p holds in the next state



$\langle \rangle p$: p holds eventually
• At least one state p holds true...



$\Box p$: p holds from now on



$p \text{ U } q$: p holds until q holds



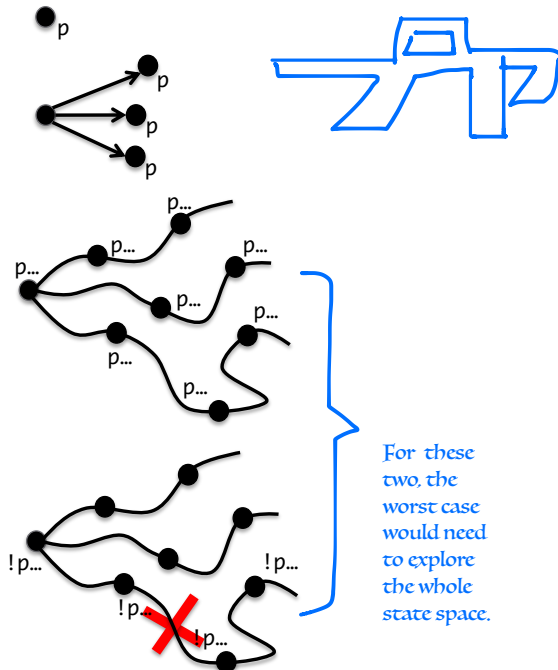
Simple algorithms can check LTL properties on a product automaton

□

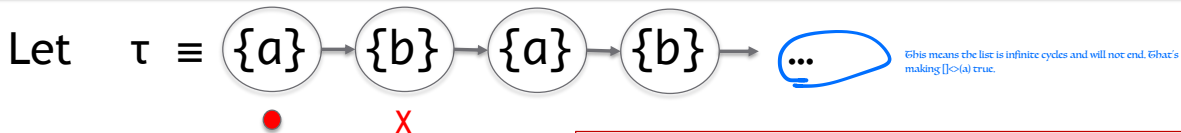
- p check p on the current/initial state
- Xp check p holds in all successors of current state

$\Box p$ find all reachable states from the current state and ensure p holds in all of them

$\langle \rangle p$ look for a path from the current state on which p is false in every state; if no such path is found, then on every path reachable from the current state, p must be true somewhere, therefore, $\langle \rangle p$ must be true

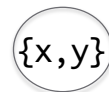


LTL Examples (for a single path)



- ✓ • $\neg b$ b is not true in the initial state, so satisfied
- ✓ • $a \rightarrow Xb$
- ✓ • $c \rightarrow a$ The precondition is not satisfied. So this implication is automatically satisfied regardless of post condition
- ✓ • $\Box(a \rightarrow Xb)$
- ✓ • $\Box(b \rightarrow Xa)$
- ✓ • $\Diamond(a)$ Eventually a : at least one a will appear in current or future state
- ✓ • $\Box\Diamond(a)$ Always eventually a : for every state, there will be at least one a that will appear in current or future state
- ✓ • $\Box(a \rightarrow \Diamond b)$
- ✓ • $a \cup b$

- a, b, c are state propositions
- τ is a path (trace/trail)
- $\{a\}, \{b\}, \dots$ indicate the state propositions that are true for each state on a path



means x, y are true, and nothing else in this state

Caution: In LTL “ \rightarrow ” is implication; not a statement separator like it is in Promela

LTL is defined in the global scope: Normally only global variables of a Promela model can appear in an LTL formula to define state propositions

If you want to refer to a local variable....

```
bit gx;  
  
ltl p {  $\Box((gx == 0) \rightarrow \Diamond(gx == 1))$  }  
  
active proctype X() {  
    bit x;  
  
    d_step { x = 0; gx = x; }  
    do  
        :: x == 1; d_step { x--; gx = x; }  
        :: x == 0; d_step { x++; gx = x; }  
    od;  
}
```

can't refer to x here

Define a global var gx
that tracks x, and copy
x to gx whenever x is
updated

See “Remote References” slides in Appendix for ways of accessing local variables or local symbolic states in a model or in an LTL



Example: Sema4.pml

```
mtype { P, V }  
chan s = [0] of { mtype }  
  
byte count = 0; // counting users in critical section  
bool wantIt[2] = 0; // user i wants the semaphore  
bool getIt[2] = 0; // user i got the semaphore  
  
active proctype Sema() {  
  do  
    :: s!P; s?V;  
  od  
}  
  
proctype User(byte i) {  
  do  
    :: wantIt[i] = 1;  
    s?P ->  
    count++;  
    wantIt[i] = 0; getIt[i] = 1; // critical section  
    count--;  
    s!V; getIt[i] = 0  
  od;  
}  
  
init {  
  run User(0);  
  run User(1);  
}
```

- Let's define some *observable global variables* for the simplified version of the Semaphore example

- Define the mutex property as an LTL
- Define a progress property for each user as an LTL:
 - always, if a user process wants the semaphore, it eventually gets the semaphore

Note that observation variables are not used to control the behavior in any way!



Sema4.pml

```
mtype { P, V }
chan s = [0] of { mtype }

byte count = 0; // counting users in critical section
bool wantIt[2] = 0; // user wants the semaphore
bool getIt[2] = 0; // user got the semaphore

active proctype Sema() {
  do
    :: s!P; s?V;
  od
}

proctype User(byte i) {
  do
    :: wantIt[i] = 1;
    s?P ->
    count++;
    wantIt[i] = 0; getIt[i] = 1;
    count--;
    s!V; getIt[i] = 0
  od;
}

init {
  run User(0);
  run User(1);
}

/*
  Verification -> Liveness, Acceptance cycles, Apply never claim
  Specify name of ltl
*/
ltl mutex { [](count < 2) }
ltl wantItGetIt0 { [](wantIt[0] -> <>getIt[0]) }
ltl wantItGetIt1 { [](wantIt[1] -> <>getIt[1]) }
```

'You want it, you get it.'

- *Verification -> Liveness*
 - *Acceptance cycles*
 - *Apply never claim*
 - *Turn off Use partial order reduction (Advanced Options)*
 - *Specify in-model LTL formula*
 - mutex
 - wantItGetIt0
 - wantItGetIt1
- Which ones fail, which ones pass?
 - $\text{Sema4} \models \text{mutex}$?
 - $\text{Sema4} \models \text{wantItGetIt0}$?
 - $\text{Sema4} \models \text{wantItGetIt1}$?
- Check console output
- Check failure trace by running a guided simulation
- Notice <<<<START OF CYCLE>>>>



Sema4.pml

Verification result:

pan: ltl formula wantItGetIt1

pan:1: **acceptance cycle** (at depth 11)

pan: wrote Sema4.pml.trail

State-vector 60 byte, depth reached 30, errors: 1

14 states, stored

0 states, matched

14 transitions (= stored+matched)

0 atomic steps

hash conflicts: 0 (resolved)

Stats on memory usage (in Megabytes):

0.001 equivalent memory usage for states (stored*(State-vector + 1))

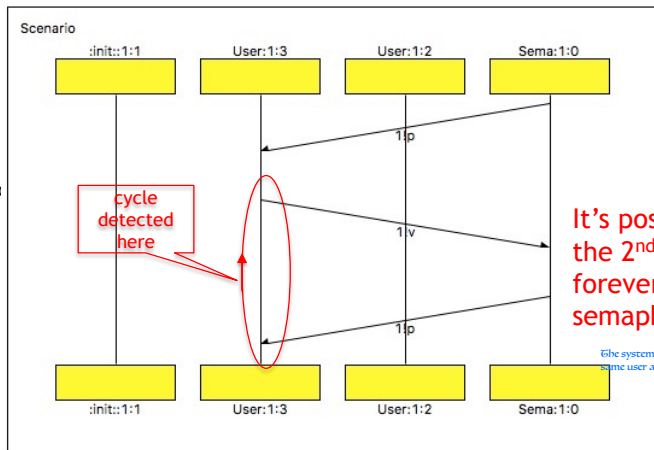
0.287 actual memory usage for states

128.000 memory used for hash table (-w24)

0.534 memory used for DFS stack (-m10000)

128.730 total actual memory usage

Sema4 \neq wantItGetIt1



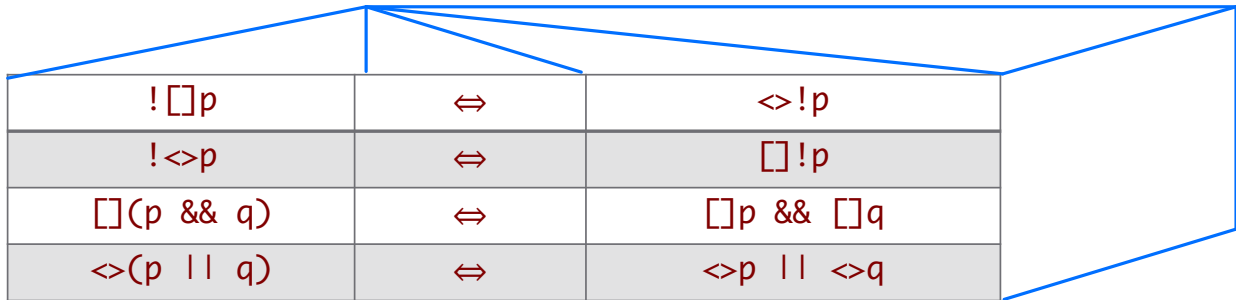
Frequently used LTL formulae

S: state
P: path

Formula	Pronounced	Interpretation	Type
$\Box p$	always p	invariance	S
$\Diamond p$	eventually p	guarantee	P
$p \rightarrow \Diamond q$	p implies eventually q	response	P
$\Box \Diamond p$	always eventually p (p holds <i>infinitely often</i>)	recurrence (progress)	P
$\Box (p \rightarrow \Diamond q)$	always, p implies eventually q	recurrent response	P
$\Diamond \Box p$	eventually always p (p becomes <i>persistent</i>)	stability or non-progress	P
$\Box q \text{ U } (p \ \&\& \ !q)$	p (strictly) precedes q (if q ever happens)	(strict) precedence	P
$\Diamond p \rightarrow \Diamond q$ p may before q, q may before p	eventually p implies eventually q	correlation	P

If p is true for the first time, q must not hold true at the same time.

Basic LTL equivalence rules



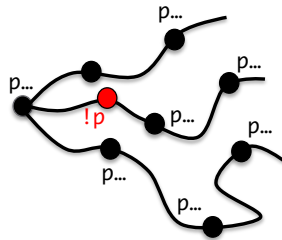
$\neg \Box p$	\Leftrightarrow	$\Diamond \neg p$
$\neg \Diamond p$	\Leftrightarrow	$\Box \neg p$
$\Box (p \ \&\& \ q)$	\Leftrightarrow	$\Box p \ \&\& \ \Box q$
$\Diamond (p \ \ q)$	\Leftrightarrow	$\Diamond p \ \ \Diamond q$

False friends

$\neg \Box p$	\nleftrightarrow	$\Box \neg p$
$\neg \Diamond p$	\nleftrightarrow	$\Diamond \neg p$

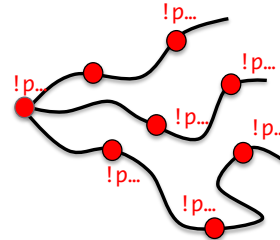
Some false LTL rules

$!\Box p$	\Leftrightarrow	$\Box !p$
$!\Diamond p$	\Leftrightarrow	$\Diamond !p$



$!\Box p$

\leftarrow



$\Box !p$

Extreme Caution!

\models : satisfies

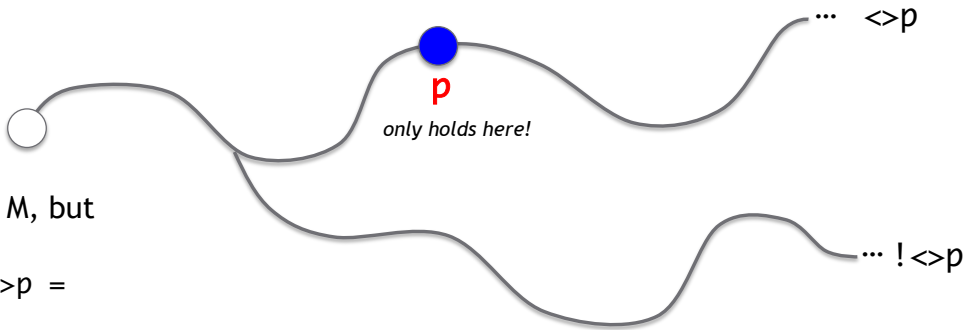
$M \not\models \varphi$ does not imply $M \models \neg\varphi$
(it's possible that neither holds)

Reason: implicit universal quantification

Suppose $\varphi = \langle \rangle p$ such that

p holds somewhere in some trace of M , but
it holds nowhere in other traces

Then neither $\langle \rangle p$ nor its negation $\neg \langle \rangle p =$
 $\Box \neg p$ will hold for M





Exercises

Try to work this out in a group of 2-3: 10 mins

Express each requirement as an LTL formula based on these system state propositions

- reqS: true if request sent; false otherwise
- reqR: true if request received; false otherwise
- ackR: true if acknowledgement received; false otherwise

Requirements

- the request is eventually sent, and once sent, it is always eventually received
 - after receiving an acknowledgement, another acknowledgement cannot be received until after a new request has been sent
-

Exercises: (non-unique) solutions

because we don't know the whole context



- reqS: request sent
 - reqR: request received
 - ackR: acknowledgement received
-
- the request is eventually sent, and once sent, it is always eventually received

(\leftrightarrow reqS)

Exercises: (non-unique) solutions



- reqS: request sent
 - reqR: request received
 - ackR: acknowledgement received
-
- the request is eventually sent, and once sent, it is always eventually received
 - the request is eventually sent, and always, once sent, it is always eventually

$(\langle \rangle \text{reqS}) \ \&\& \ (\Box(\text{reqS} \rightarrow \langle \rangle \text{reqR}))$ // recurrent response

Exercises: (non-unique) solutions



- reqS: request sent
- reqR: request received
- ackR: acknowledgement received
- the request is eventually sent, and once sent, it is always eventually received

$(\langle \rangle \text{reqS}) \ \&\& \ (\Box(\text{reqS} \rightarrow \langle \rangle \text{reqR}))$ // recurrent response

- after receiving an acknowledgement, another acknowledgement cannot be received until after a new request has been sent
ackR \rightarrow $X(!\text{ackR} \ U \ \dots)$

Exercises: (non-unique) solutions



- reqS: request sent
- reqR: request received
- ackR: acknowledgement received
- the request is eventually sent, and once sent, it is always eventually received

$(\langle \rangle \text{reqS}) \ \&\& \ (\Box(\text{reqS} \rightarrow \langle \rangle \text{reqR}))$ // recurrent response

- after receiving an acknowledgement, another acknowledgement cannot be received until after a new request has been sent

$\text{ackR} \rightarrow \neg(\neg \text{ackR} \cup (\text{reqS} \ \&\& \ \neg \text{ackR}))$ // strict precedence

Exercises: (non-unique) solutions



- reqS: request sent
 - reqR: request received
 - ackR: acknowledgement received
-
- the request is eventually sent, and once sent, it is always eventually received

$(\langle \rangle \text{reqS}) \ \&\& \ (\Box(\text{reqS} \rightarrow \langle \rangle \text{reqR}))$ // recurrent response

- after receiving an acknowledgement, another acknowledgement cannot be received until after a new request has been sent (... decide that this must be true always)

$\Box(\text{ackR} \rightarrow X(\neg \text{ackR} \cup (\text{reqS} \ \&\& \ \neg \text{ackR})))$ // strict precedence

START WORKING ON L4

Review the instructions, due dates, submission limits, and partnering rules both on Canvas and on Vocareum carefully before starting!

If you activate Vocareum solo without specifying a partner, you will need to finish L4 alone!

Appendix

Promela Quick Reference

<http://spinroot.com/spin/Man/Quick.html>

Example: checking for pure atomicity

Suppose we want to check that none of the atomic clauses in our model are ever blocked (i.e. pure atomicity).

1. Add a global bit variable:

```
bit aflag;
```



2. Change all atomic clauses to:

```
atomic {  
  stat1;  
  aflag=1;  
  stat2  
  
  ...  
  
  statn  
  aflag=0;  
}
```



3. Check that aflag is always 0.

```
[!]!aflag
```

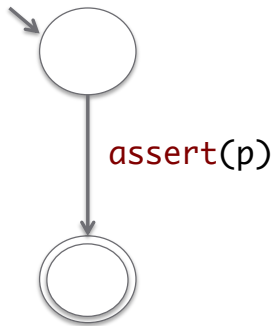
e.g.

```
active process monitor {  
  assert(!aflag);  
}
```

Global invariants can be defined in different ways, some more efficient than others

```
active proctype monitor()  
{  
    assert(p) ;  
}
```

```
active proctype monitor()  
{  
    do  
        :: assert(p)  
    od  
}
```



Same effect, but more efficient since it has a single state



More on deadlocks...

- Automatically detected by default in Spin in most verification modes (can be turned off - don't check *Invalid end states* in SpinRCP *Verification*)
 - Whenever Spin reaches a global state with no enabled transitions and that are not valid **end states**, it will report an error
 - The PC location just before the closing curly brace of a **proctype** is always a valid **end state** for that process (will not lead to reporting of a deadlock)
 - **timeout** statements may disable deadlock states
 - *be careful about not masking possible deadlock behavior by a global timeout!*
-

More on end states: **end**-state labels define legal termination

If a process counter (PC) location is labeled as an **end** state, a state ending at that location will be treated as an intentional and valid end state (deadlock will not be reported even if execution is not possible out of that state)

End-state labels are prefixed with the keyword **end** and must be unique in each process

- E.g., **end**, **end1**, **endstate**, **end**-state



End-state example: Disjktra's semaphore: an infinitely running system (never reaches an end state)

```
mtype { P, V }
```

```
chan s = [0] of { mtype }
```

```
active proctype Sema() {
  bit free = 1;
  do
    :: (free) -> s!P; free = 0;
    :: (!free) -> s?V; free = 1;
  od;
}
```

rendez-vous channel
ensures atomicity of
semaphore operations,
which must be indivisible

```
active [3] proctype user() {
  do
    :: s?P; // enter
    critical: skip;
    s!V; // leave
  od;
}
```

3 users competing for the
same semaphore

Since the users keep
requesting the semaphore,
no invalid end states found
in Spin verification

Run *Random Simulation*, see that P and V
never gets interleaved and the system
runs indefinitely

More on rendez-vous: rendez-vous reduces interleaving

SemaRVOk.pml

```
#define MAX 0 // this works
#define USERS 3

mtype { P, V }
byte count = 0;

chan s = [MAX] of { mtype }

active proctype Sema() {
  do
    :: s?P; s?V;
  od;
}

active [USERS] proctype User() {
  do
    :: s!P; // enter
       critical: count++;
       leaving: count--;
       s!V; // leave
  od;
}

active proctype mutex () {
  assert(count <= 1);
}
```

SemaRVOk \models mutex

SemaRegNotOk.pml

```
#define MAX 2 // this doesn't work
#define USERS 3

mtype { P, V }
byte count = 0;

chan s = [MAX] of { mtype }

active proctype Sema() {
  do
    :: s?P; s?V;
  od;
}

active [USERS] proctype User() {
  do
    :: s!P; // enter
       critical: count++;
       leaving: count--;
       s!V; // leave
  od;
}

active proctype mutex () {
  assert(count <= 1);
}
```

SemaRegNotOk $\not\models$ mutex

Verification -> Safety -> Assertion violations | Simulation -> Random

More on rendez-vous: rendez-vous reduces interleaving

SemaRVOk.pml

```
#define MAX 0 // this works
#define USERS 3

mtype { P, V }
byte count = 0;

chan s = [MAX] of { mtype }

active proctype Sema() {
  do
    :: s?P; s?V;
  od;
}

active [USERS] proctype User() {
  do
    :: s!P; // enter
       critical: count++;
       leaving: count--;
       s!V; // leave
  od;
}

active proctype mutex () {
  assert(count <= 1);
}
```

SemaRVOk \models mutex

SemaRegNotOk.pml

```
#define MAX 2 // this doesn't work
#define USERS 3

mtype { P, V }
byte count = 0;

chan s = [MAX] of { mtype }

active proctype Sema() {
  do
    :: s?P; s?V;
  od;
}

active [USERS] proctype User() {
  do
    :: s!P; // enter
       critical: count++;
       leaving: count--;
       s!V; // leave
  od;
}

active proctype mutex () {
  assert(count <= 1);
}
```

S & V ops no
longer
indivisible

SemaRegNotOk $\not\models$ mutex

Verification -> Safety -> Assertion violations | Simulation -> Random

Rendez-vous reduces interleaving

SemaRVOk.pml

```
#define MAX 0 /* this works */
#define USERS 3

mtype { p, v };
byte count = 0;

chan s = [MAX] of { mtype };

active proctype Sema()
{
    do
        :: s?p; s?v;
    od
}

active [USERS] proctype User()
{
    do
        :: s!p;          // enter
           critical: count++;
           leaving: count--;
           s!v;          // leave
    od
}

active proctype mutex () {
    assert(count <= 1)
}
```

SemaRVOk \models mutex

SemaRegNotOk.pml

```
#define MAX 2 /* this doesn't work */
#define USERS 3

mtype { p, v };
byte count = 0;

chan s = [MAX] of { mtype };

active proctype Sema()
{
    do
        :: s?p; s?v;
    od
}

active [USERS] proctype User()
{
    do
        :: s!p;          // enter
           critical: count++;
           leaving: count--;
           s!v;          // leave
    od
}

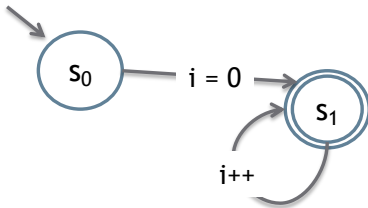
active proctype mutex () {
    assert(count <= 1)
}
```

SemaRegNotOk $\not\models$ mutex

Verification -> Safety -> Assertion violations | Simulation -> Random

A clarification on *global vs. local* state

What program code could this FSA represent

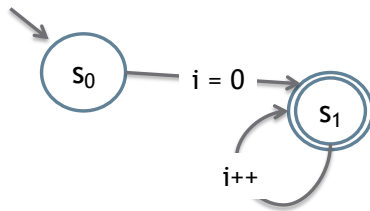


i is one-byte unsigned integer

What is the complete state space of this FSA



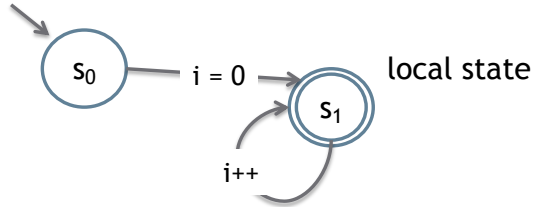
```
active proctype forever() {  
    byte i;  
    i = 0;  
s1:  
    do  
        :: i++  
    od  
}
```



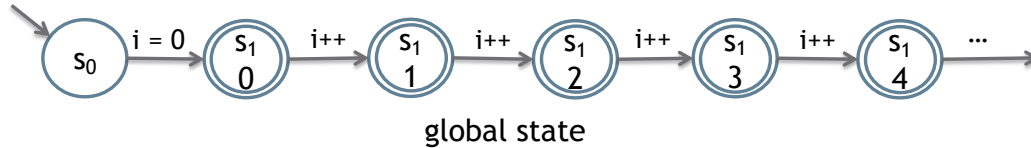
i is one byte unsigned integer

How many global states?
(in Spin)

What is the complete state space of this FSA



i is one byte unsigned integer



257 global states?

Local state is different from global state

*Process state as defined in a model is **local**: represents execution/process counter, or node-label of underlying FSA*

*System state is **global** and comprises:*

- identities of all active processes
- local state of all active processes
- values of all local variables of all active processes
- contents of all channels
- values of all global variables

LTL interpretation examples

Formula	Interpretation
$\Box p$	p is invariantly true
$\Diamond \Box !p$	p eventually becomes invariantly false
$\Box \Diamond !p$	p always eventually becomes false, at least once more
$\Box (q \rightarrow \Diamond !p)$	q always implies eventually p becomes false
$\Box (q \rightarrow !p)$	q always implies p is false in the same state

LTL operator precedence: unary over binary

Always use brackets to be prudent!

Temporal illusions

Be careful!

$p \rightarrow q$

p implies q only in initial state

$\Box(p \rightarrow q)$

p implies q in every state, locally for that state

$\Box(p \rightarrow \langle \rangle q)$

ok, but p and q may be true together (no strict ordering)

$\Box(p \rightarrow X\langle \rangle q)$

q becomes true only after p becomes true (strict ordering, *but* trivially true if p is never true)
or did you mean this?

$(\langle \rangle p) \ \&\&$

$(\Box(p \rightarrow X\langle \rangle q))$

p must eventually be true and
q becomes true after p becomes true

A word about recurrent response

$$\Box(p \rightarrow \Diamond q)$$

Problem

- May trivially hold if p is never true
- May need to check that:
 - (a) p can be true in all traces (eventually p) - *may be too strong*
 - (b) p can be true along some trace

Solution

- (a) check pre-condition separately by checking $\Diamond p$ holds
- (b) check pre-condition separately by $\Box !p$, but for failure
 - If $\Box !p$ fails, then the counterexample is proof that p holds somewhere for some trace

Remember

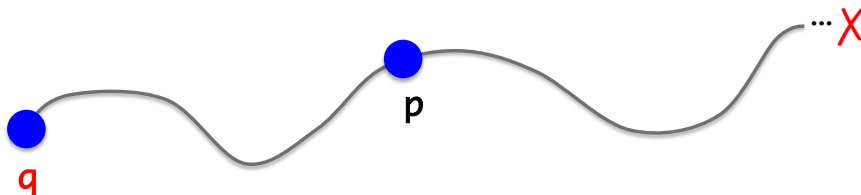
- Checking $!\Box !p = \Diamond p$ holds for a model is not the same as checking $\Box !p$ fails for that model
- Both $!\Box !p$ and $\Box !p$ may fail for a model simultaneously because of implicit universal quantification

Recurrent response

Always Eventually

↓ ↓

$\Box (p \rightarrow \Diamond q)$

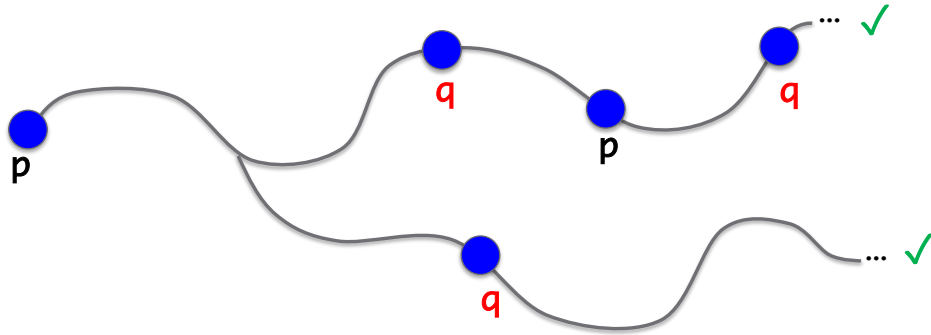


Recurrent response

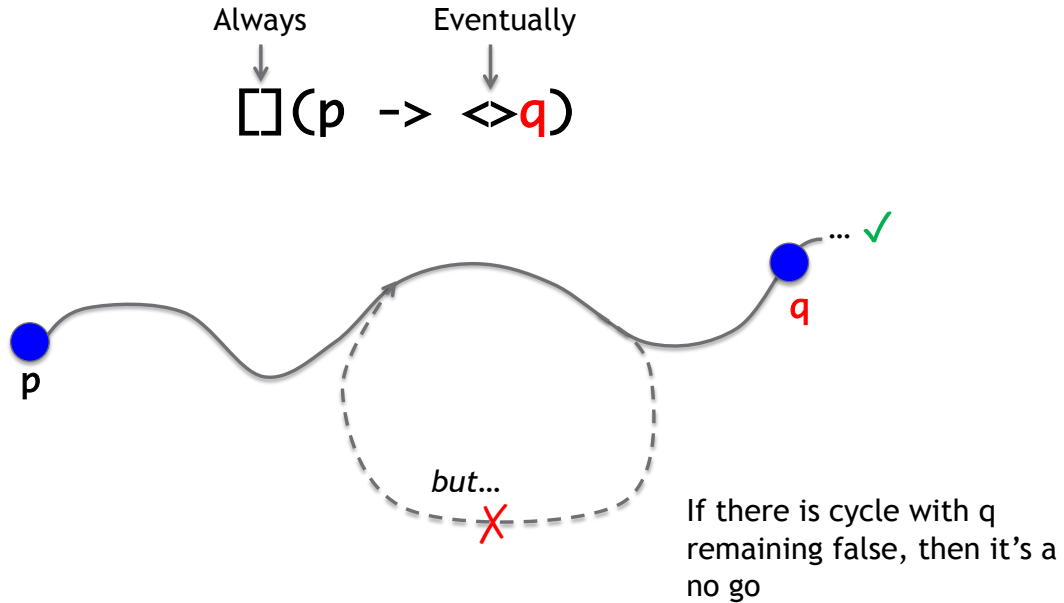
Always Eventually

↓ ↓

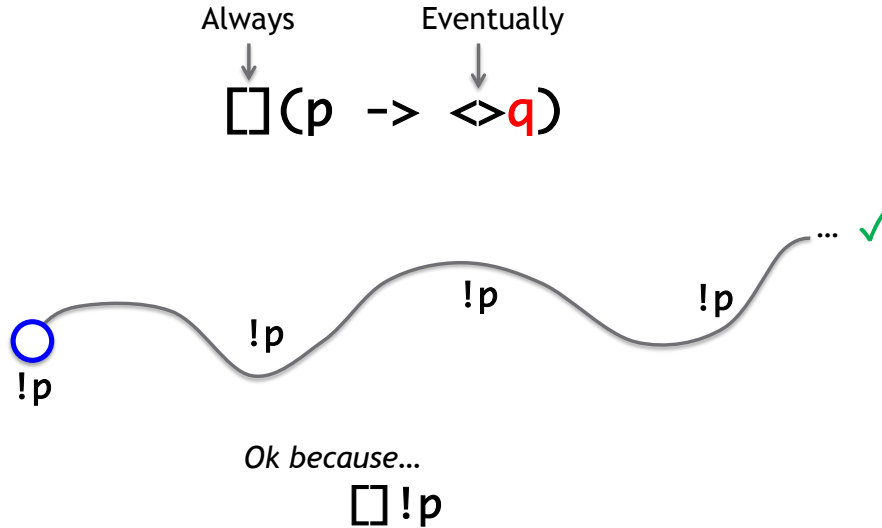
$\Box(p \rightarrow \Diamond q)$



Recurrent response



Recurrent response

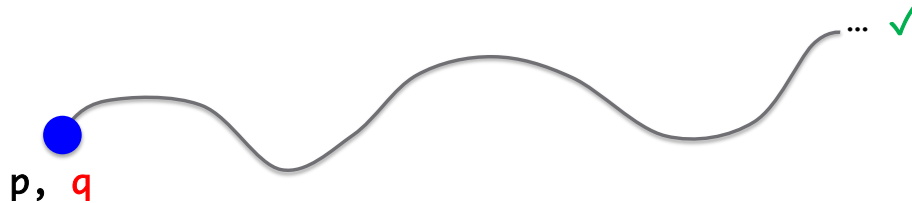


Recurrent response

Always Eventually

↓ ↓

$\Box(p \rightarrow \Diamond q)$



Safety/state properties: mechanisms

- end-state labels

```
proctype Sema()  
  endSema:  
  do  
    :: s!p;  
    :: s?v;  
  od  
}
```

- assert

- monitor processes

```
active proctype always() {  
  assert(count < 2) ;  
}
```

- LTL: exclusive use of temporal operator []

```
ltl alwaysLessThanTwo { [](count < 2) }
```

Liveness/path properties: mechanisms

- LTL: inclusion of temporal operators $\langle \rangle$, X , U

`ltl alwaysEventuallyZero { $\square \langle \rangle$ (count == 2) }`

- progress labels (not covered)
- never claims (not covered)

Liveness properties -- Progress: the system will move to a productive state

`ltl willProgress { <>(mutex > 0) }`

```
bit x, y;  
byte mutex = 0;  
  
active proctype A() {  
    x = 1;  
    y == 0;  
    mutex++;  
    // A in critical section  
    mutex--;  
    x = 0;  
}
```

```
active proctype B() {  
    y = 1;  
    x == 0;  
    mutex++;  
    // B in critical section  
    mutex--;  
    y = 0;  
}
```


Livness properties - Termination: the system will eventually stop at a valid end state

```
bool done = false;

ltl stops { <>(done) }

active proctype Q() {
    int i = 0;
    do
        :: i < MAX -> i++;
        :: i < MAX -> break;
        :: i == MAX -> break;
    od
    done = true;
}
```



Specifying termination with LTL (liveness)

```
bool ok = false;

ltl p { <>(ok) } // fails

active proctype Q() {
    byte i = 0;
    byte MAX = 10;
    do
        :: i < MAX -> i++;
        :: i <= MAX -> i--;
        :: i == MAX -> break;
    od
    ok = true;
}
```

Verification result:

warning: only one claim defined, -N ignored
pan:1: acceptance cycle (at depth 18)
pan: wrote Termination.pml.trail

(Spin Version 6.3.2 -- 17 May 2014)
Warning: Search not completed
+ Partial Order Reduction

Full statespace search for:
never claim + (p)
assertion violations + (if within scope of
claim)
acceptance cycles + (fairness enabled)
invalid end states - (disabled by never
claim)

State-vector 28 byte, depth reached 24, errors: 1
13 states, stored (16 visited)
1 states, matched
17 transitions (= visited+matched)
0 atomic steps



AltBit4.pml

```
mtype = { msg, ack }
chan to_sndr = [2] of { mtype, bit }
chan to_rcvr = [2] of { mtype, bit }
bit msgSent[2] = 0; // for LTL
bit ackReceived[2] = 0; // for LTL

active proctype Sender() {
  bit seq_out = 0;
  bit seq_in = 0;
  // obtain first message
  do
  :: if
    :: to_rcvr!msg(seq_out);
    // send message
    :: skip; // or loose it
  fi;
  msgSent[seq_out] = 1;
  msgSent[seq_out] = 0;
  to_sndr?ack(seq_in);
  ackReceived[seq_in] = 1;
  ackReceived[seq_in] = 0;
  if
  :: seq_in == seq_out ->
    // obtain new message
    seq_out = 1 - seq_out;
  :: else; // retransmit same message
  fi;
od;
}
```

```
ltl msg0Sent { <>msgSent[0] }
ltl msg1Sent { <>msgSent[1] }
ltl recover0 { [] (msgSent[0] -> <>ackReceived[0]) }
ltl recover1 { [] (msgSent[1] -> <>ackReceived[1]) }
```

```
active proctype Receiver() {
  bit seq_in = 1; // important
  do
  :: if
    :: to_rcvr?msg(seq_in); // receive msg
    :: timeout; // recover from msg loss
  fi;
  if
  :: to_sndr!ack(seq_in); // send ack
  :: skip; // or loose it
  fi;
od;
}
```

- Verification of each LTL
 - Liveness
 - Acceptance cycles
 - Apply never claim
 - Turn off Use partial order reduction (Advanced Options)
- Which ones fail, which ones pass?
- Check console output
- Check failure trace



AltBit4.pml

```
mtype = { msg, ack }
chan to_sndr = [2] of { mtype, bit }
chan to_rcvr = [2] of { mtype, bit }
bit msgSent[2] = 0; // for LTL
bit ackReceived[2] = 0; // for LTL

ltl msg0Sent { <>msgSent[0] }
ltl msg1Sent { <>msgSent[1] }
ltl recover0 { [] (msgSent[0] -> <>ackReceived[0]) }
ltl recover1 { [] (msgSent[1] -> <>ackReceived[1]) }

active proctype Sender() {
    bit seq_out = 0;
    bit seq_in = 0;
    // obtain first message
    do
        :: if
            :: to_rcvr!msg(seq_out); // send message
            :: skip; // or loose it
        fi;
        msgSent[seq_out] = 1;
        msgSent[seq_out] = 0;
        to_sndr?ack(seq_in);
        ackReceived[seq_in] = 1;
        ackReceived[seq_in] = 0;
        if
            :: seq_in == seq_out ->
                // obtain new message
                seq_out = 1 - seq_out;
            :: else; // retransmit same message
        fi;
    od;
}

active proctype Receiver() {
    bit seq_in = 1; // important
    do
        :: if
            :: to_rcvr?msg(seq_in); // receive msg
            :: timeout; // recover from msg loss
        fi;
        if
            :: to_sndr!ack(seq_in); // send ack
            :: skip; // or loose it
        fi;
    od;
}
```

- *Verification for each LTL*
 - *Liveness*
 - *Acceptance cycles*
 - *Apply never claim*
- $\text{AltBit4} \models \text{msg0Sent}$
- $\text{AltBit4} \not\models \text{msg1Sent}$
- $\text{AltBit4} \not\models \text{recover0}$
- $\text{AltBit4} \not\models \text{recover1}$
- Must limit message losses



AltBit5.pml

Limit losses!

```
active proctype Sender() {
    bit seq_out = 0;
    bit seq_in = 0;
    // obtain first message
    do
    :: if
    :: (lostMsg > MAXLOST) -> atomic {
        to_rcvr!msg(seq_out);
        lostMsg = 0;
    } // definitely send message
    :: (lostMsg <= MAXLOST) ->
        if
        :: atomic { to_rcvr!msg(seq_out);
            lostMsg = 0;
        } // send message
        :: lostMsg++; // or loose it
        fi;
    fi;
    msgSent[seq_out] = 1;
    msgSent[seq_out] = 0;
    to_sndr?ack(seq_in);
    ackReceived[seq_in] = 1;
    ackReceived[seq_in] = 0;
    if
    :: seq_in == seq_out ->
        // obtain new message
        seq_out = 1 - seq_out;
    :: else; // retransmit same message
    fi;
    od;
}
```

```
...
byte lostMsg = 0; // count lost msg
byte lostAck = 0; // count lost ack
byte MAXLOST = 3; // limit max no. of losses
```

```
ltl recover0 { [] (msgSent[0] -> <>ackReceived[0]) }
ltl recover1 { [] (msgSent[1] -> <>ackReceived[1]) }
ltl maxLoss { [] (lostMsg <= MAXLOST && lostAck <= MAXLOST) }
// these last two should all fail - both channels are lossy up to MAXLOST
ltl FAIL_noLoss { [] (lostMsg == 0 && lostAck == 0) }
ltl FAIL_noLossMax { [] (lostMsg < MAXLOST && lostAck < MAXLOST) }
```

- AltBit5 \models recover0
- AltBit5 \models recover1
- AltBit5 \models maxLoss
- AltBit5 $\not\models$ FAIL_noLoss
- AltBit5 $\not\models$ FAIL_noLossMax

```
active proctype Receiver() {
    bit seq_in = 1; // important
    do
    :: if
    :: to_rcvr?msg(seq_in); // receive msg
    :: timeout; // recover from msg loss
    fi;
    if
    :: (lostAck > MAXLOST) -> atomic {
        to_sndr!ack(seq_in);
        lostAck = 0;
    } // send ack
    :: (lostAck <= MAXLOST) ->
        if
        :: atomic { to_sndr!ack(seq_in) ->
            lostAck = 0;
        }
        :: lostAck++ // or loose it
        fi;
    fi;
    od;
}
```

Liveness properties -- Non-starvation: the system will never indefinitely deny access to a resource (it will be fair to all users of resource)

```
bool accessed[2], reserved[2];  
#define reserveResource(i) ...  
#define useResource(i) ...
```

```
ltl guaranteeAccess0 { [] (reserved[0] -> <>accessed[0]) }  
ltl guaranteeAccess1 { [] (reserved[1] -> <>accessed[1]) }
```

```
active[2] proctype Client(int i) {  
    accessed[i] = false; requested[i] = false;  
    do // repeatedly access resource  
    :: reserveResource(i); reserved[i] = true;  
       useResource(i); accessed[i] = true;  
       reserved[i] = false; accessed[i] = false;  
    od;  
}  
...
```

Liveness properties (e.g., non-starvation) may rely on **fairness** assumptions

Weak form: If a statement becomes enabled and stays enabled, it will eventually be executed

- *if a process **waits** at a state in which it can make progress, it eventually will*

Strong form: If a statement is infinitely often enabled, it will eventually be executed

- *if a process constantly **revisits** a state in which it can make progress, it eventually will*

*We can force Spin to assume **weak fairness** when checking liveness/path properties*

Some more LTL rules

$\neg \Box p$	\Leftrightarrow	$\Diamond \neg p$
$\neg \Diamond p$	\Leftrightarrow	$\Box \neg p$
$\Box (p \ \&\& \ q)$	\Leftrightarrow	$\Box p \ \&\& \ \Box q$
$\Diamond (p \ \ q)$	\Leftrightarrow	$\Diamond p \ \ \Diamond q$
$p \ U \ (q \ \ r)$	\Leftrightarrow	$(p \ U \ q) \ \ (p \ U \ r)$
$(p \ \&\& \ q) \ U \ r$	\Leftrightarrow	$(p \ U \ r) \ \&\& \ (q \ U \ r)$
$\Box \Diamond (p \ \ q)$	\Leftrightarrow	$(\Box \Diamond p) \ \ (\Box \Diamond q)$
$\Diamond \Box (p \ \&\& \ q)$	\Leftrightarrow	$(\Diamond \Box p) \ \&\& \ (\Diamond \Box q)$

Remote references: states

`process_name[process_index]@label`

returns true (1) if the process instance is currently at the local state labeled `label`

`[process_index]` can be omitted if there is a single instance of the process

Can be used in LTL: see `StateLTL.pml`

Remote references: local variables

`process_name[process_index]@var`

returns value of local variable `var` for a process instance

`[process_index]` can be omitted if there is a single instance of the process

Can be used in LTL