

# ANNOUNCEMENTS

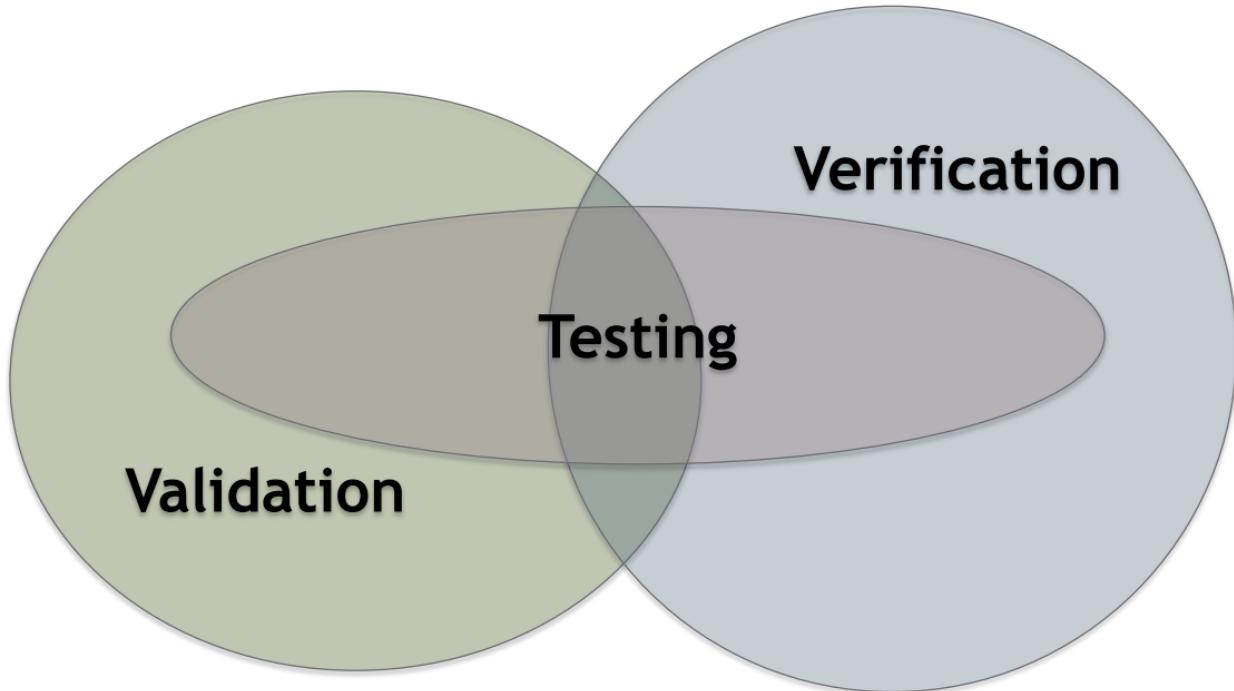
## SUCCESSION OF PREPS, READINGS, ASSIGNMENTS (CHECK CANVAS)

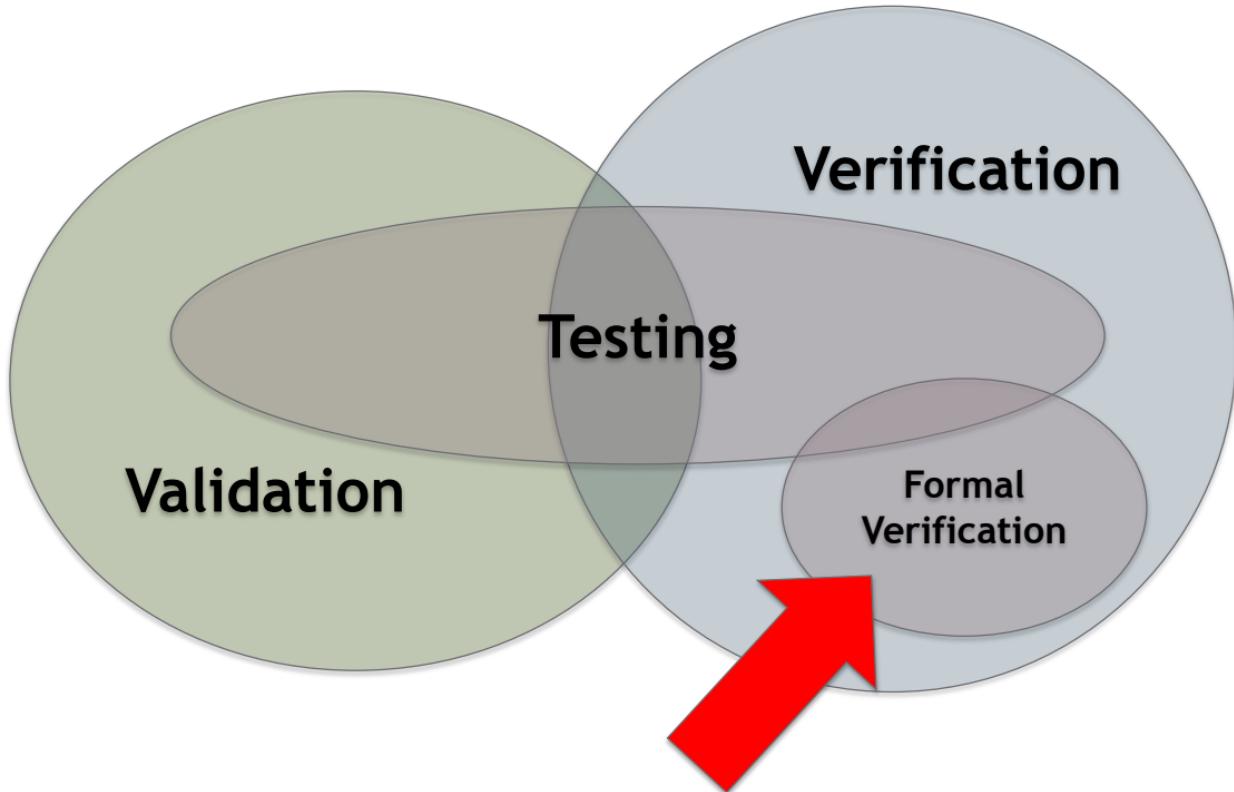
- This Thu: tool installation (Promela/Spin Prep) for class exercises and 2 articles (deadline for report for articles is Fri)
  - Tue 11/14: 2 more articles, class exercises
  - Thu 11/16: More lass exercises
  - L4: Warmup for Promela/Spin Assignment (take-home): due date on Canvas
  - A2: Promela/Spin Assignment (last assignment): due date on Canvas
- 
- FINAL EXAM: TUE DEC 12, 2:30PM PT (ROOM TBD)

*formal = opposite of ad-hoc*

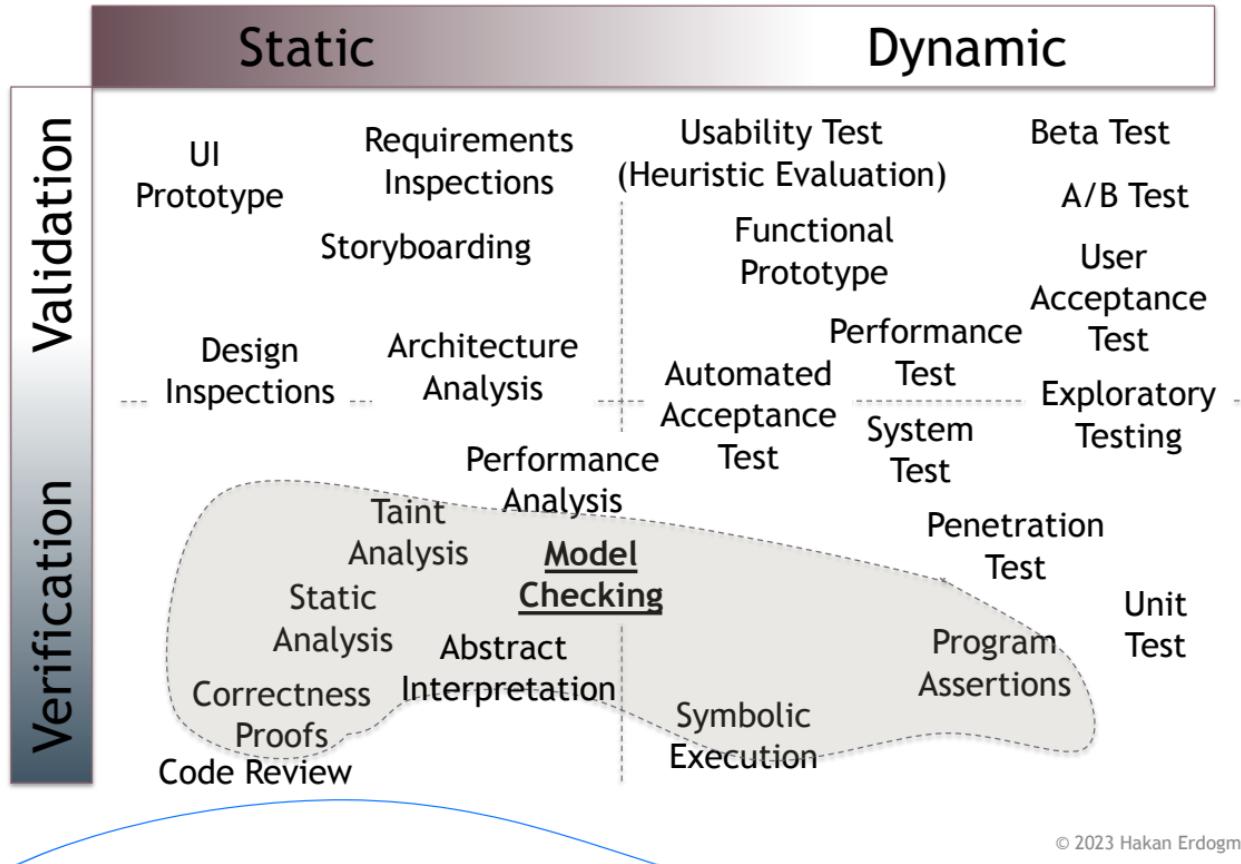
## FORMAL VERIFICATION

rigorous, algorithmic, logic-, and mathematically-based verification techniques





# V&V Landscape - *in this course*

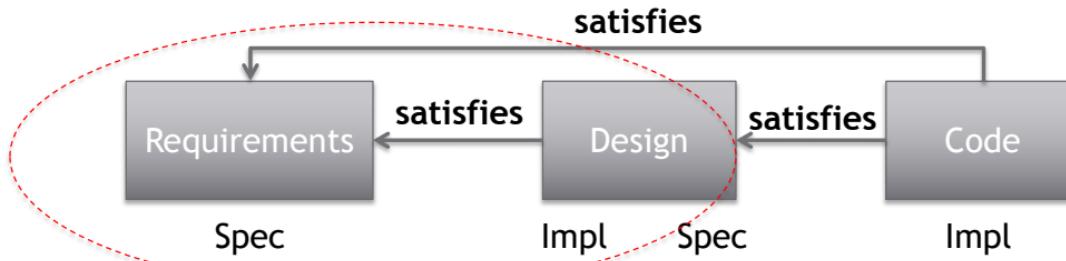


Formally: Verification is checking consistency of an *implementation* with respect to a *specification*

---

*Impl satisfies Spec*

*Implementation and specification are roles,  
not fixed artifacts*



**Verification:** checking consistency of a *model* with respect to a set of correctness *properties*

Impl

Spec

## *Model satisfies Property*

$$M \vDash P$$

Model is an abstraction of the implementation that includes only the details relevant to the property being verified (**approximation principle**)

*When:* Spec = Set of Properties = Set of Requirements

$$M \models \{ P_1, P_2, \dots, P_n \}$$

# Main Principles of V&V still apply here

# Main Principles of V&V still apply here

- **Partitioning:** divide and conquer
- **Approximation:** making the problem easier
- **Visibility:** making information accessible
- **Feedback:** providing actionable information
- **Repeatability:** better to fail every time than sometimes
- **Redundancy:** triangulating results by using different techniques

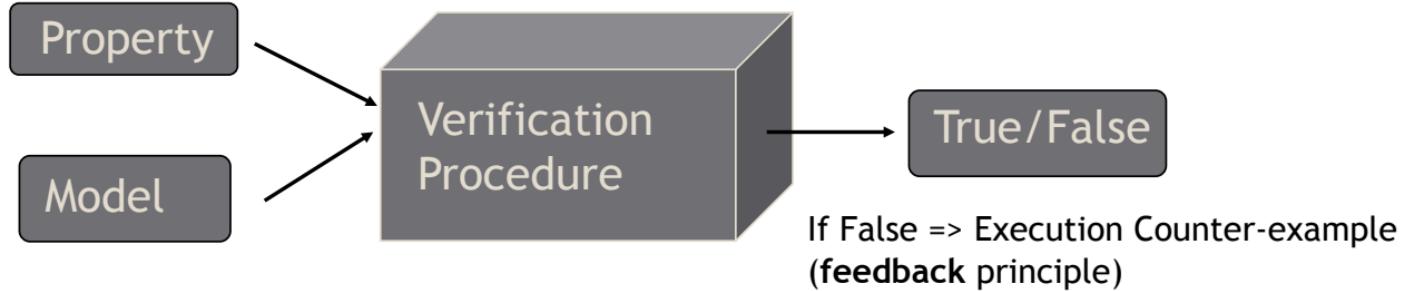
## Examples of properties

---



- The system always shuts down when the pressure exceed 100psi
  - The system never refuses inputs
  - A request is always followed by an acknowledgement if it's accepted
  - Every message sent will eventually be received
  - The requests will be processed in the order they are sent
  - After pressing the call button, the elevator will eventually go to the calling floor unless the cancel button is pressed
-

~~ever~~  
YOU CAN'T ~~ALWAYS~~ GET WHAT YOU WANT



PROVING CORRECTNESS PROPERTIES IS UNDECIDABLE IN  
ITS FULL GENERALITY...

*Restrict to subset of relevant and useful  
properties on finite-state models*

# Also important: How realistic/faithful is the model?

---

*Depends on the property*

- If all details relevant to the property being verified are included in the model, then realistic/faithful (*sound*)
- Unrealistic/unfaithful to the extent of missing or changed detail relevant to the property being verified (*unsound*)

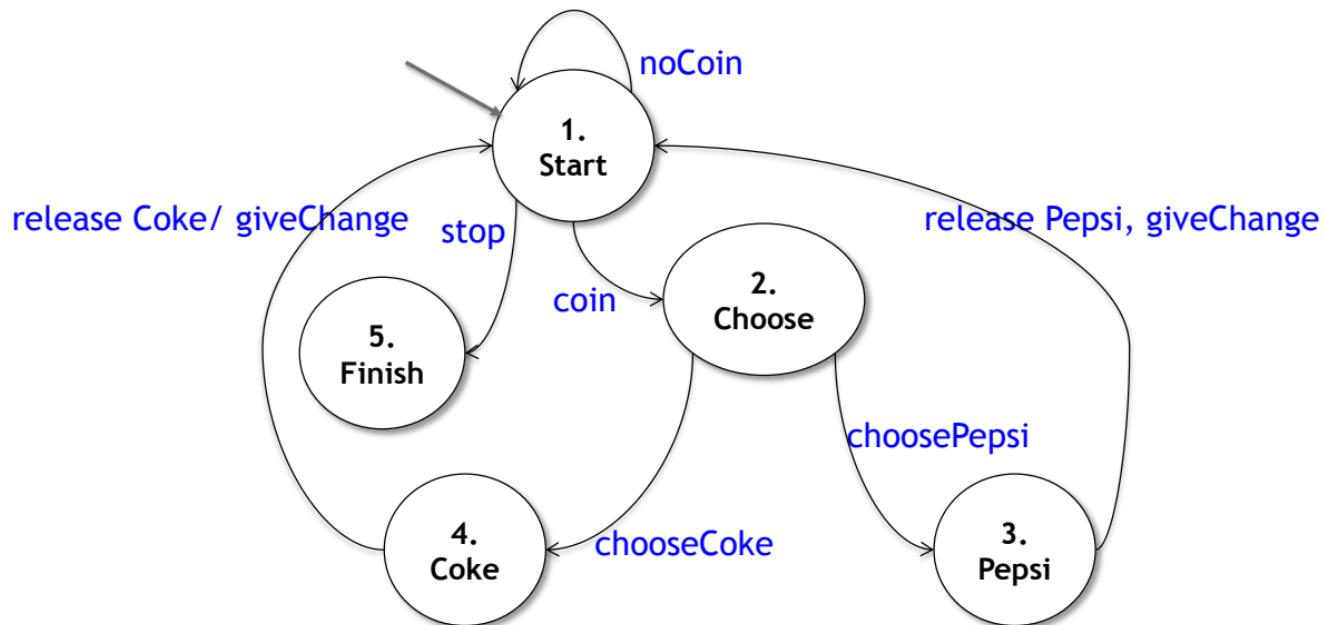
# Building models

---

- What do we need to know to build a model?
- Consider a simple vending machine
  - A customer inserts coins, selects a beverage and receives a can of soda
  - Basically, a state machine
    - Starts in initial state and changes based on user interaction
- We need to know
  - What *states* the vending machine can be in
  - How the machine *transitions* from one state to another
  - In which state the machine starts

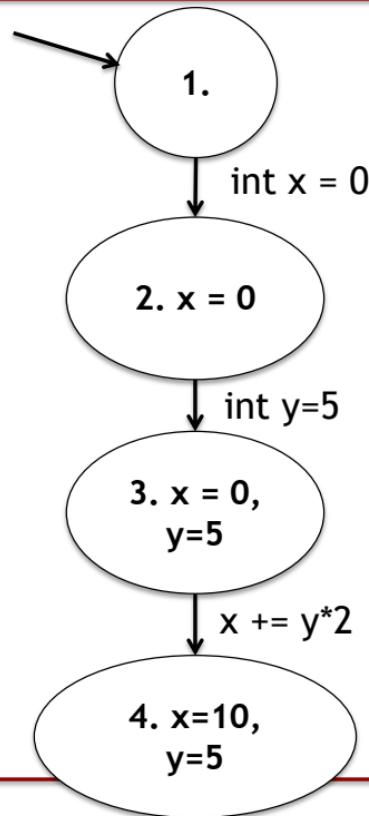


# VENDING MACHINE MODEL



# Modeling a sample program as a state machine

```
int x = 0;  
int y = 5;  
x += y*2;
```



---

**Formal verification techniques that use models, like static analysis, can find faults that testing is not good at finding!**

# Which types of faults testing is good at finding?

---



- Faults that we can predictably reproduce at execution
  - Often related to computations and data flow in sequential and deterministic systems

# Which types of faults is testing poor at finding?

---



Faults that we cannot predictably reproduce at execution

- Often related to complex, unpredictable **control flow**
  - when there is inherent *nondeterminism*
- Common in **concurrent systems**
  - parallel programs
  - distributed systems

This is where formal verification techniques come into play

## Concurrent systems are *composed of multiple communicating components*

---

- each component has own thread of execution
  - components communicate or synchronize through different mechanisms
    - message passing
      - events
      - remote procedure call (RPC)
    - request-response protocol (HTTP)
    - shared data
    - *rendez-vous (handshake)*
  - no guarantees with respect to relative speed of execution
-



---

## Examples of concurrent systems



# Examples of concurrent software systems

---

- Multi-threaded programs
  - Actor systems / Agent systems
  - Communicating processes
  - Network applications
  - Operating systems
  - Embedded systems
  - Sensor networks / IoT systems
  - Distributed web applications
  - Client-server applications
  - Grid applications
  - Communication protocols
- ...

# Most software is concurrent

---

*... at some sufficiently low-enough or sufficiently high-enough level of abstraction*

Whether we treat (model) the system under verification as *concurrent* or *sequential* depends on the properties of interest...



---

What are some common concurrency-related faults?

# Common concurrency-related faults

---



- Deadlocks
    - *Inability to make progress*
  - Race conditions
    - *Untimely or out-of-order access to shared resources*
  - Livelocks
    - *Unproductive progress ("spinning its wheels")*
  - Starvation
    - *Unfair scheduling, unfair access to resources*
-



## Think-Pair-Share: Find familiar examples of each category

---

- Deadlocks
    - *Inability to make progress*
  - Race conditions
    - *Untimely or out-of-order access to shared resources*
  - Livelocks
    - *Unproductive execution*
  - Starvation
    - *Unfair scheduling, unfair access to resources*
-

Testing cannot easily catch concurrency faults  
due to inability to control OS, communication, concurrency infrastructure

---

- *relative execution speeds of underlying components/threads/processes*
- *opaque scheduling schemes of execution environments*
- *some nondeterminism being totally invisible*

We cannot easily reproduce concurrency faults during testing!

*relevant*

- We need to consider *all possible behaviors*, independent of any invisible nondeterminism!
  - **Formal verification can do this, but we will need to model the system first!**
-

# Motivating example: mutual exclusion

Two threads that need a resource  $r$  are trying to control access to  $r$ .

Shared variables  $x$ ,  $y$  are used to put a *lock* on  $r$ .

Another shared variable  $stop$  controls termination.

Thread X

```
stop = false;  
while (!stop) {  
    x = 1; ----- reserve -----  
    while (y == 1) {} ----- busy wait -----  
    // can access r  
    dSW(r); ----- access (critical section) -----  
    x = 0; ----- release -----  
}
```

Thread Y

```
stop = false;  
while (!stop) {  
    y = 1;  
    while (x == 1) {}  
    // can access r  
    dSW(r);  
    y = 0;  
}
```

How can you write a test that makes sure only one thread accesses  $r$  at a time? (mutex property)

And that at least one thread is able to access it? (progress property)



# Motivating example: how do we formulate these properties?

---

Introduce a shared *observation* variable `mutex` (init to 0) that track how many threads are in their critical section

Thread X

```
stop = false;  
while (!stop) {  
    x = 1;  
    while (y == 1) {}  
    mutex++;  
    dSW(r);  
    mutex--;  
    x = 0;  
}
```

Thread Y

```
stop = false;  
while (!stop) {  
    y = 1;  
    while (x == 1) {}  
    mutex++;  
    dSW(r);  
    mutex--;  
    y = 0;  
}
```

- **Mutex property:** value of `mutex` should never be greater than 1
- **Progress property:** one thread can enter its critical section; *i.e.*, `mutex` will eventually become 1

---

— What now?



# Let's try to do this with testing

Introduce a shared history array  $h[ ]$  and shared count  $c$  (init to 0).  $h[c]$  stores a snapshot of variable mutex. In  $dSW(r)$ , store value of mutex in  $h[c]$  and increment  $c$ .

```
Thread X
stop = false;
while (!stop) {
    x = 1;
    while (y == 1) {}
    mutex++;
    dSW(r); h[c++] = mutex;
    mutex--;
    x = 0;
}
```

```
Thread Y
stop = false;
while (!stop) {
    y = 1;
    while (x == 1) {}
    mutex++;
    dSW(r);
    mutex--;
    y = 0;
}
```

```
Test: threadX.start(); threadY.start(); Thread.sleep(50);
stop = true; end = c; progress = false;
threadX.join(200); threadY.join(200);
for (int i = 0; i < end; i++) {
    assertTrue(h[i] <= 1); // mutex
    if (h[i] > 0) progress = true;
}
assertTrue(progress); // progress
```

Run this test 10 times. It always passes. Your strategy is working. You're done!



# Wait: something is wrong!

Change nothing. Come back next day. Re-execute the test. Now it suddenly fails? Test reveals that `c == 0` at the end. What the heck?

```
Thread X
stop = false;
while (!stop) {
    x = 1;
    while (y == 1) {}
    mutex++;
    dSW(r);
    mutex--;
    x = 0;
}
```

```
Thread Y
stop = false;
while (!stop) {
    y = 1;
    while (x == 1) {}
    mutex++;
    dSW(r);
    mutex--;
    y = 0;
}
```

```
Test: threadX.start(); threadY.start(); Thread.sleep(50);
stop = true; end = c; progress = false;
threadX.join(200); threadY.join(200);
for (int i = 0; i < end; i++) { // loop body never executed
    assertTrue(h[i] <= 1); // mutex
    if (h[i] > 0) progress = true;
}
assertTrue(progress); // progress
```

Come back tomorrow. Run the test. Now it fails! Why?



# Wait: something is wrong!

Change nothing. Come back next day. Re-execute the test. Now it suddenly fails? Test reveals that `c == 0` at the end. What the heck?

```
Thread X
stop = false;
while (!stop) {
    x = 1;
    while (y == 1) {}
    mutex++;
    dSW(r);
    mutex--;
    x = 0;
}
```

```
Thread Y
stop = false;
while (!stop) {
    y = 1;
    while (x == 1) {}
    mutex++;
    dSW(r);
    mutex--;
    y = 0;
}
```

```
Test: threadX.start(); threadY.start(); Thread.sleep(50);
stop = true; end = c; progress = false;
threadX.join(200); threadY.join(200);
for (int i = 0; i < end; i++) { // loop body never executed
    assertTrue(h[i] <= 1); // mutex
    if (h[i] > 0) progress = true;
}
assertTrue(progress); // progress
```

Come back tomorrow. Run the test. Now it fails! Why?



# Java Mutex Example - Generalized Version

---

For a generalized version of the previous example with N clients/threads and observe its unpredictability, try:

<https://github.com/cmusv-svvt/Mutex-Java>



# Let's try a different approach

---

- Model the system as communicating finite state automata
- Search the global state space
- See if there are any undesirable states:
  - that violate the *mutex* property
  - that violate the *progress* property

Thread X

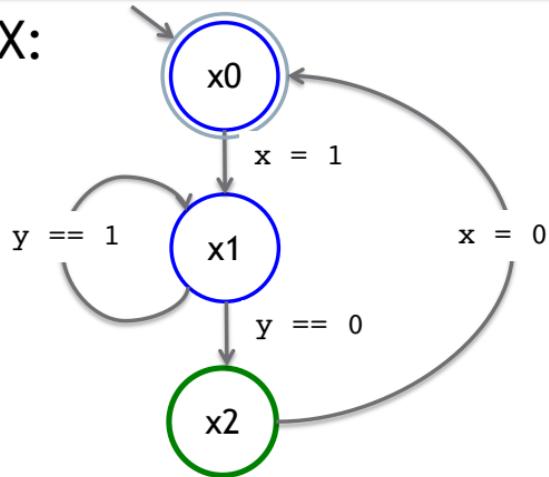
```
stop = false;
while (!stop) {
    x = 1;
    while (y == 1) {}
    mutex++;
    dSW(r);
    mutex--;
    x = 0;
}
```

Thread Y

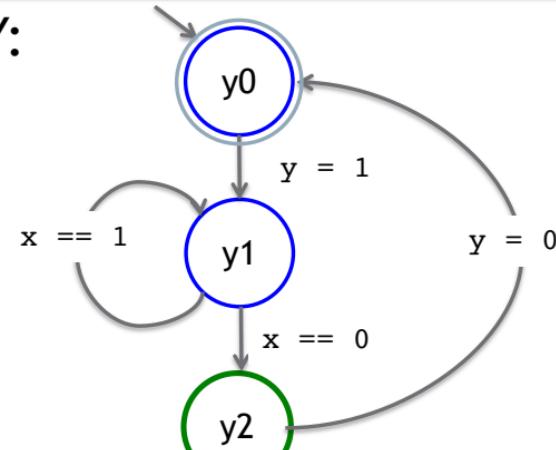
```
stop = false;
while (!stop) {
    y = 1;
    while (x == 1) {}
    mutex++;
    dSW(r);
    mutex--;
    y = 0;
}
```

# Finite State Automata for X and Y

X:



Y:



```
while (true) {
    x = 1;
    while (y == 1) {}
    dSW(r);
    x = 0;
}
```

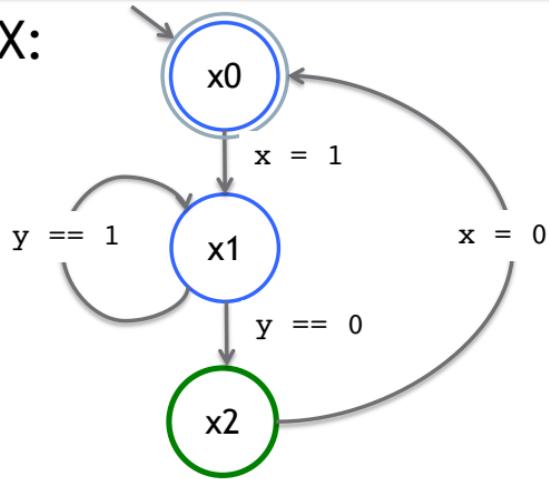
```
while (true) {
    y = 1;
    while (x == 1) {}
    dSW(r);
    y = 0;
}
```

in critical section  
outside critical section

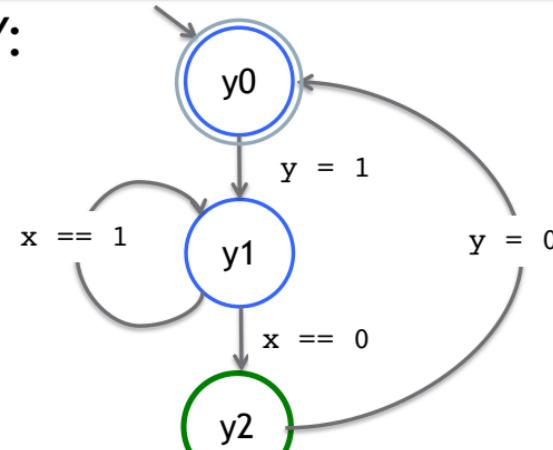


# Finite State Automata for X and Y

X:



Y:



```
while (true) {
    x = 1;
    while (y == 1) {}
    x = 0;
}
```

```
while (true) {
    y = 1;
    while (x == 1) {}
    y = 0;
}
```

in critical section

outside critical  
section

What happened to stop, mutex++, mutex--, dSW(r)?

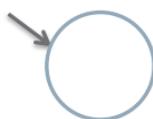


# Finite State Automata

---

- States and transitions (both finite)

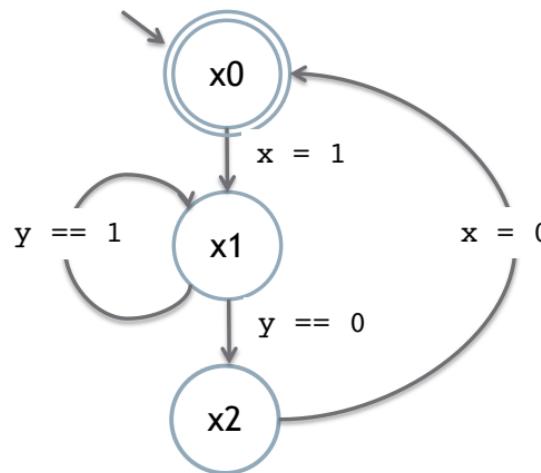
- Start state



- A set of end states



- Transitions are labeled
  - *condition*: enabled if evaluates to true
  - *action (command)* (e.g., assignment): always enabled

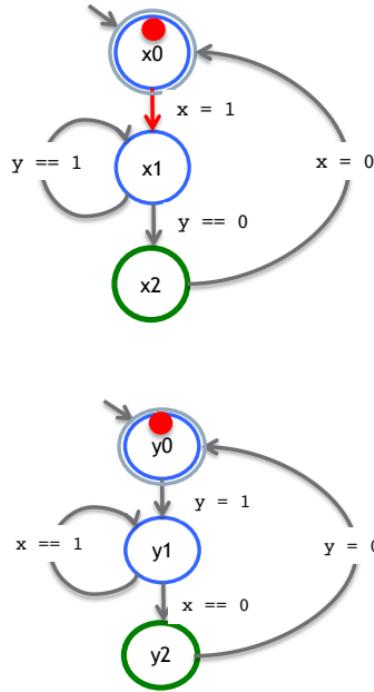
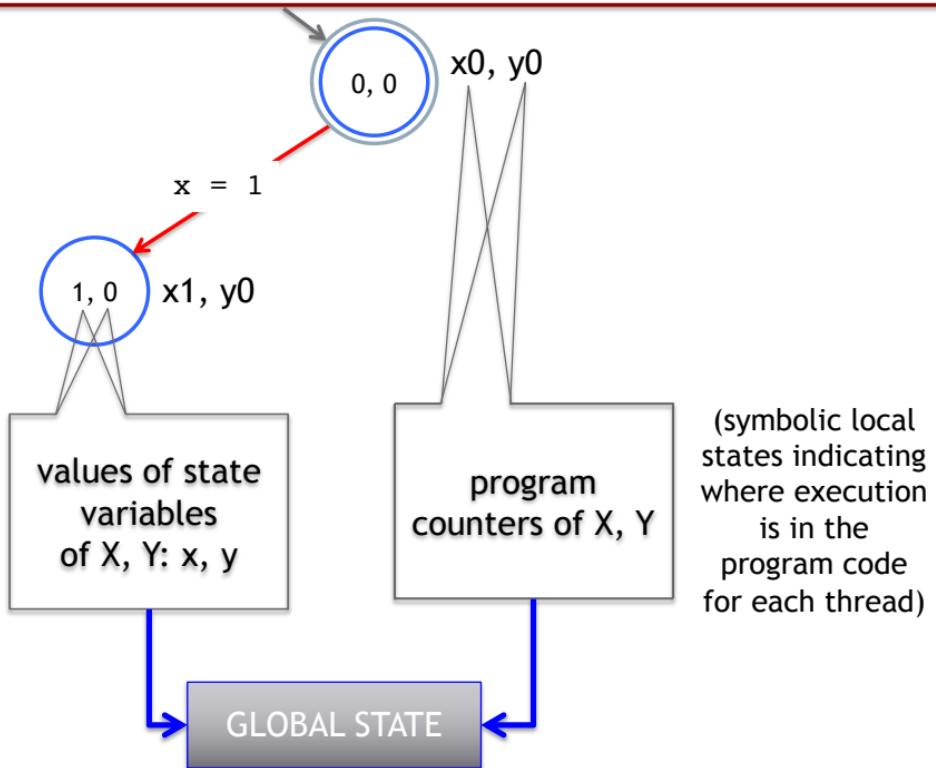


## The product FSA: $X \otimes Y$

---

- We want to search the combined (global) state space
- One FSA can execute an enabled transition at a time (emulating a single-core CPU multiplexing different threads)
- In case of multiple enabled transitions, one is randomly chosen
- We can do a depth-first or breadth-first search to visit all global states
- When we visit an already visited global state, we backtrack and look for another enabled transition to take

# Creating $X \otimes Y$ (1)

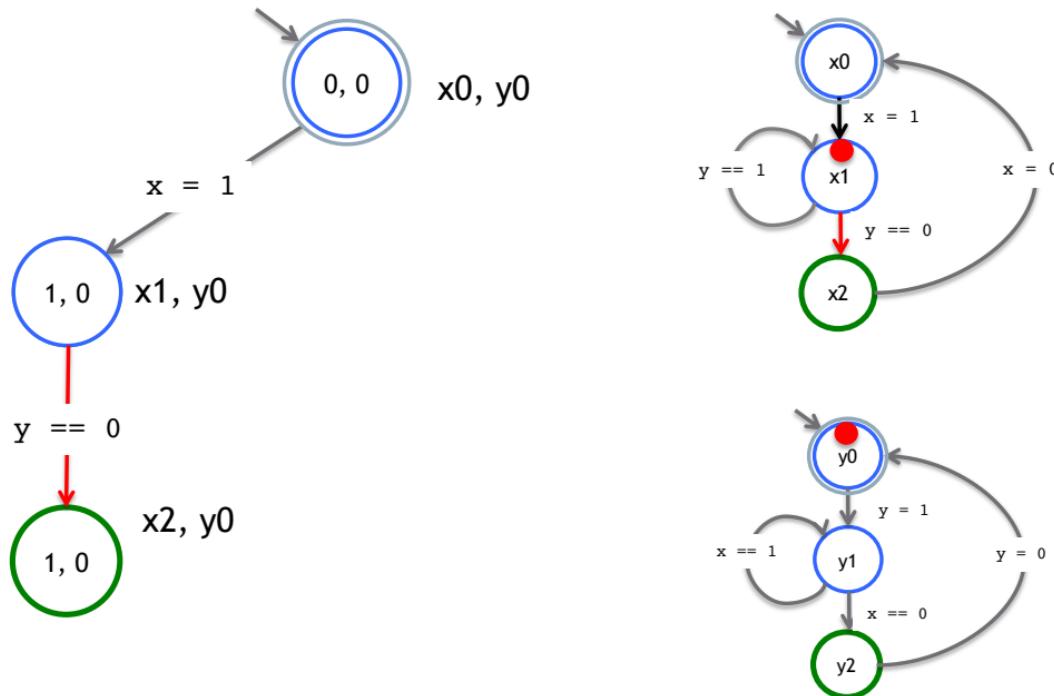


*every combination is a distinct global state*

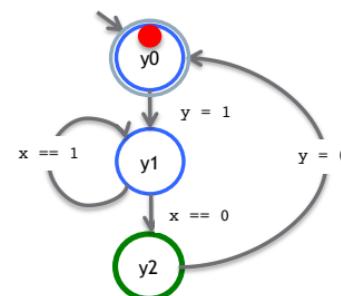
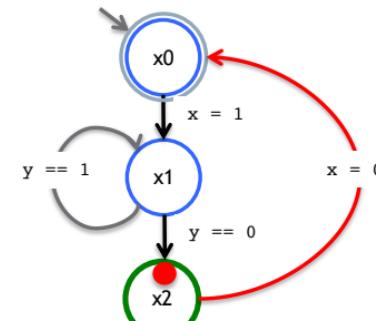
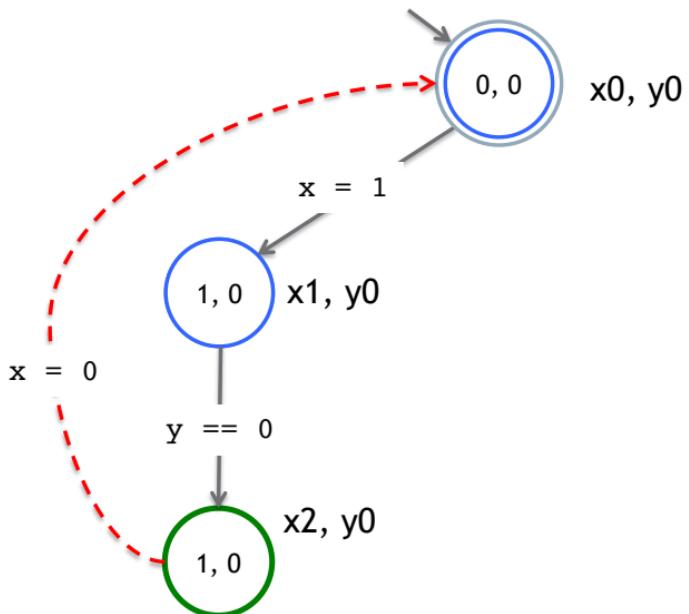


# Creating $X \otimes Y$ (2)

---



# Creating $X \otimes Y$ (3)

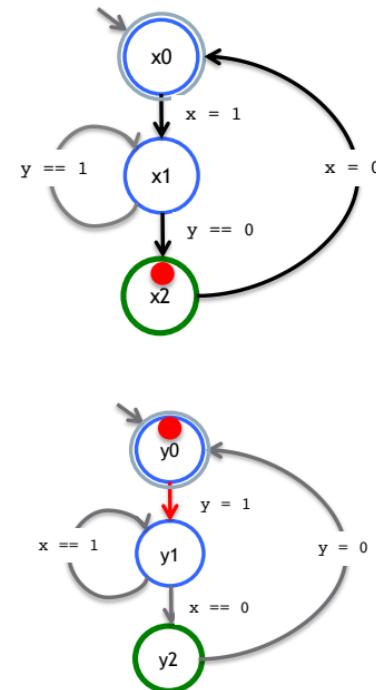
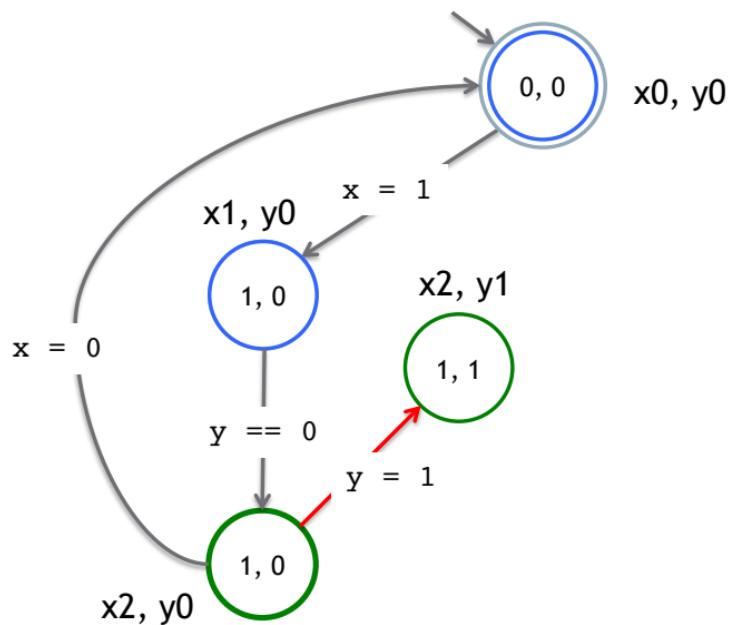


Hit previously visited state: stop and backtrack to  $x_2, y_0$

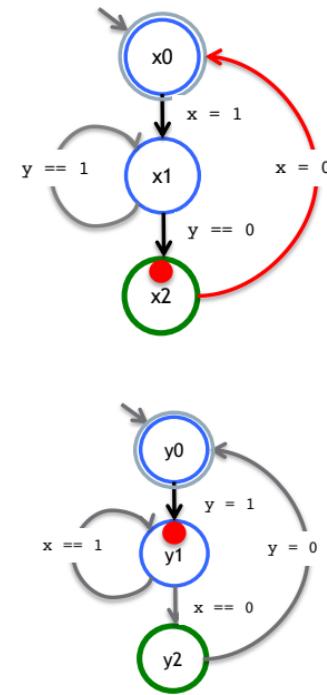
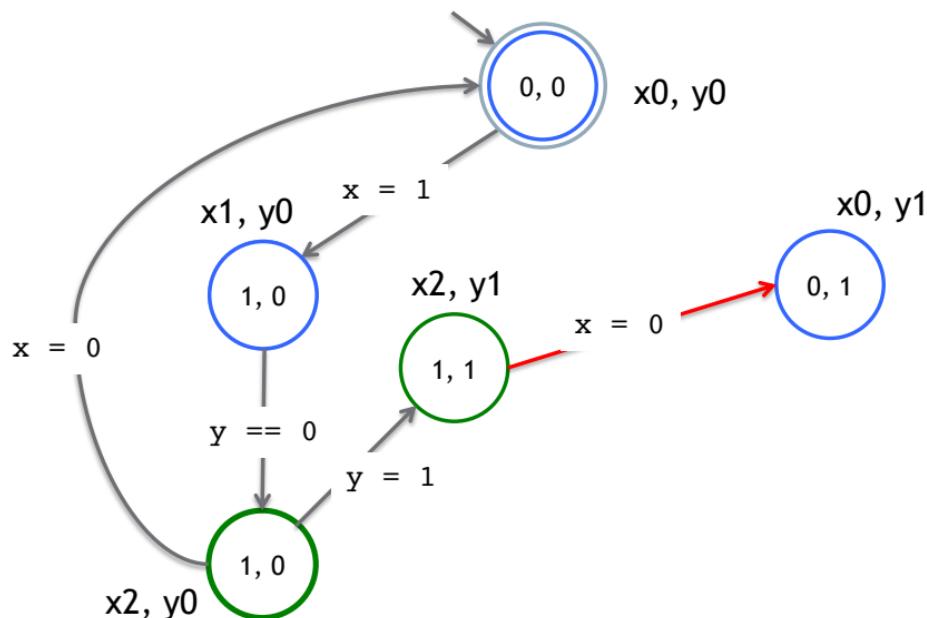


# Creating $X \otimes Y$ (4)

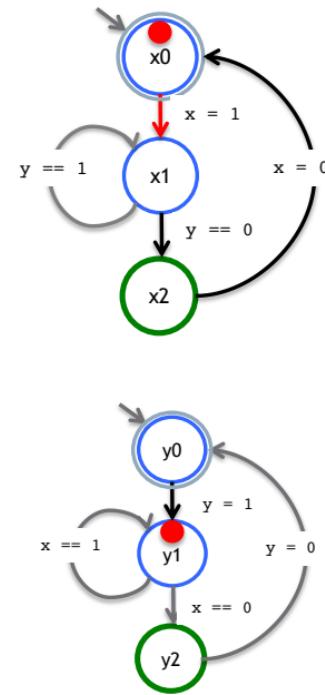
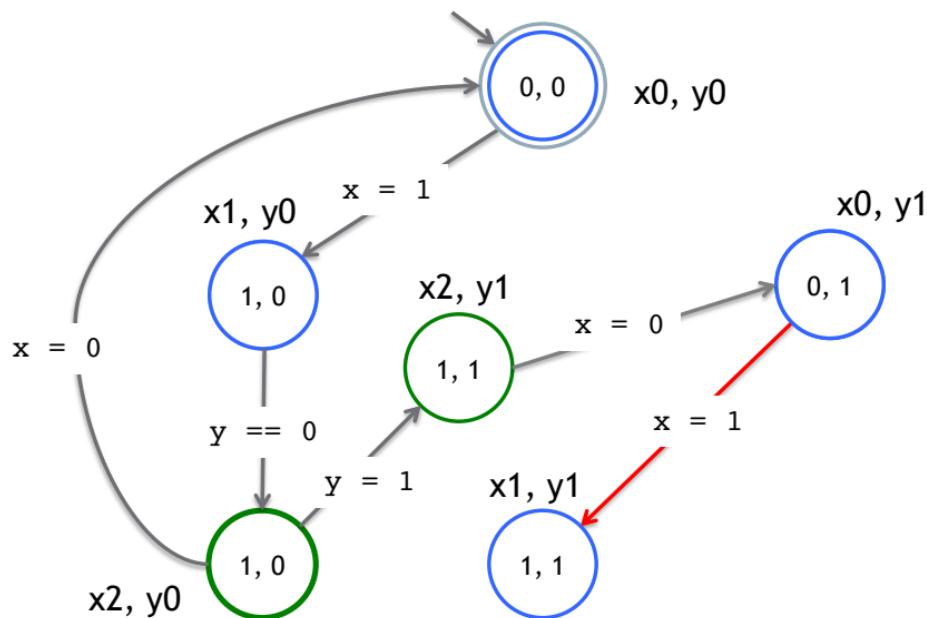
---



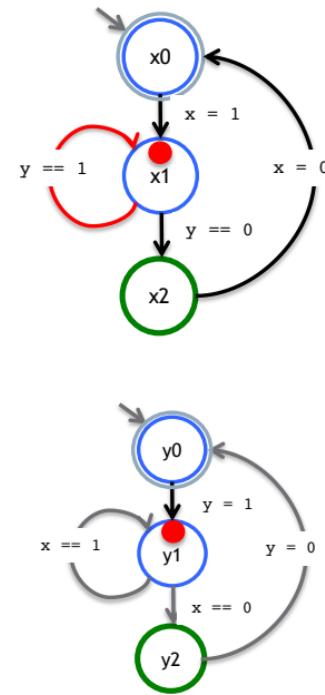
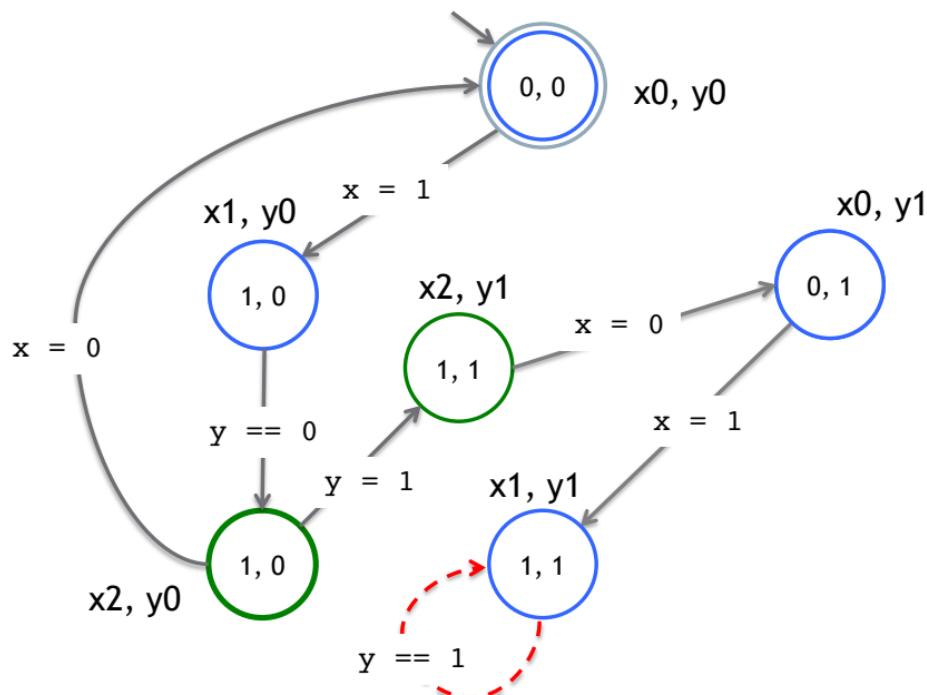
# Creating $X \otimes Y$ (5)



# Creating $X \otimes Y$ (6)



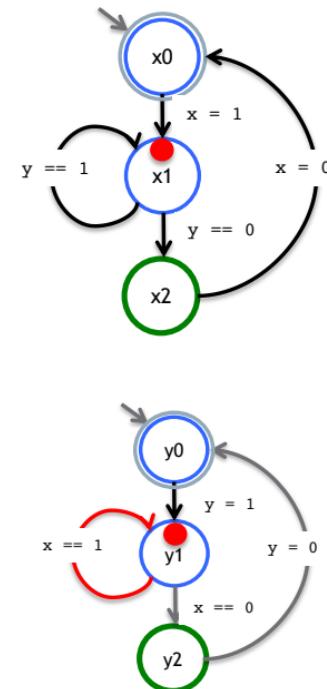
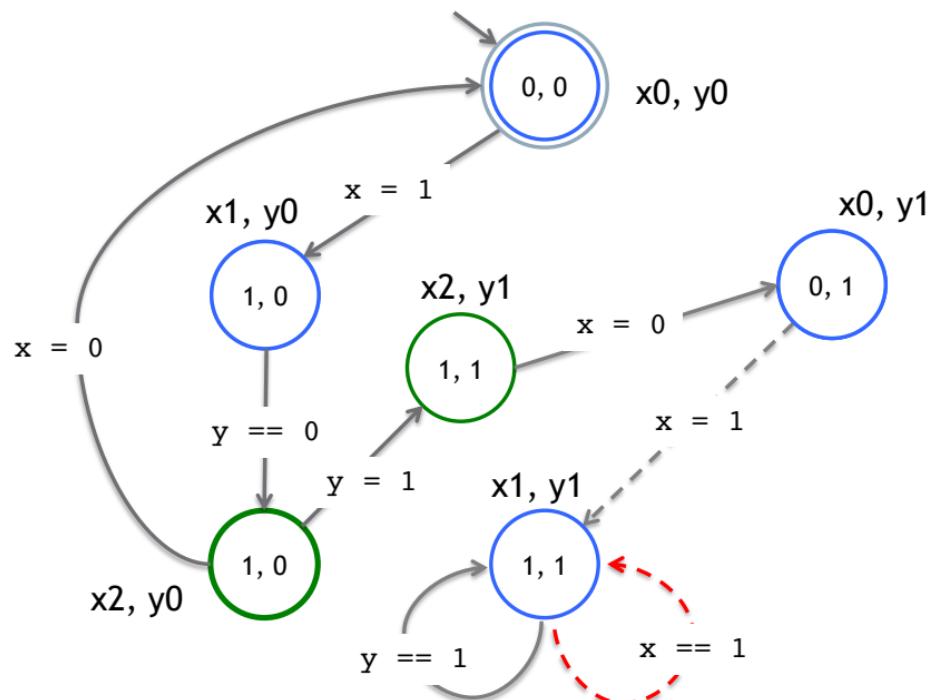
# Creating $X \otimes Y$ (7)



Hit previously visited state: stop and backtrack to  $x_1, y_1$



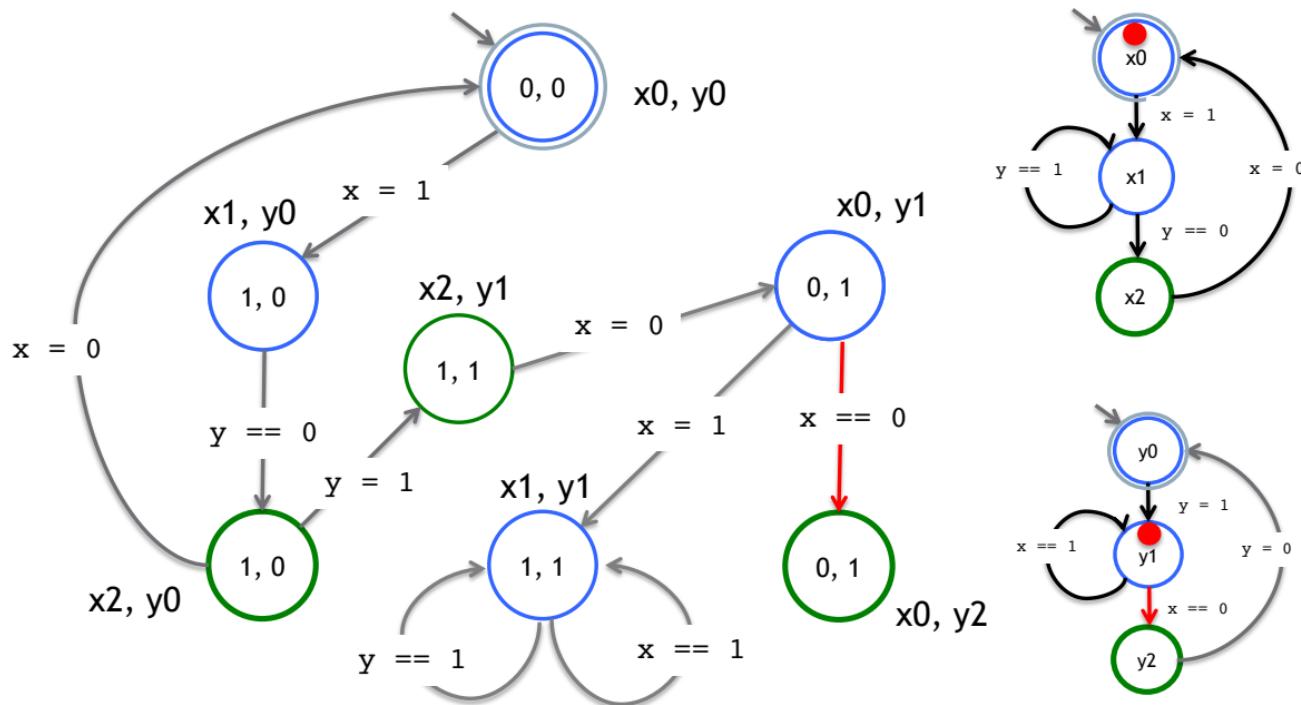
# Creating $X \otimes Y$ (8)



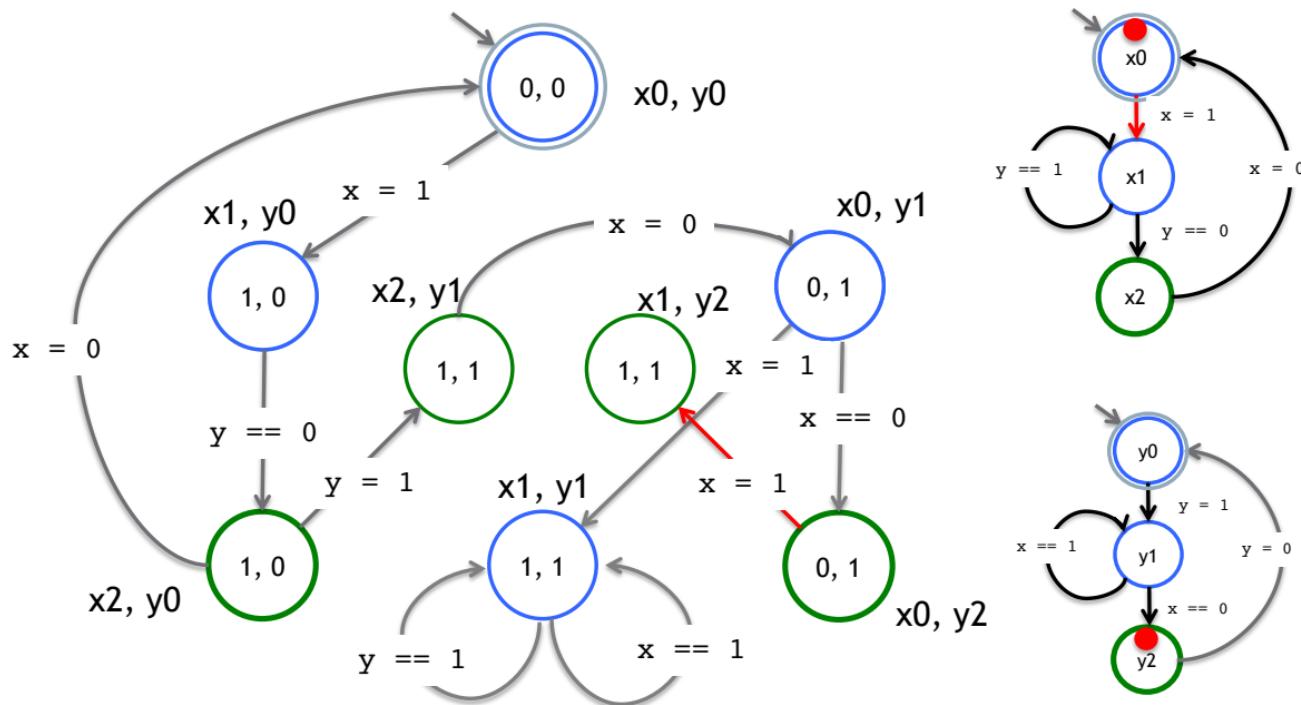
No further transitions to explore at  $x_1, y_1$ :  
backtrack to  $x_0, y_1$



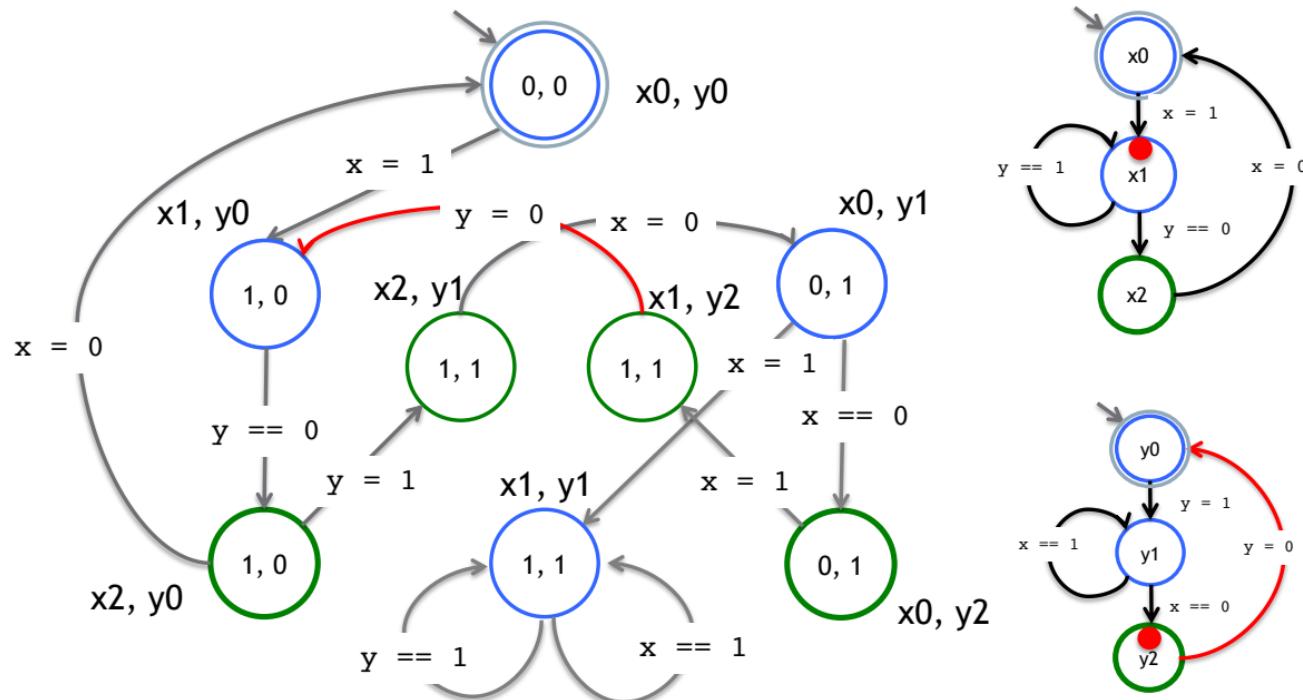
## Creating X ⊗ Y (9)



# Creating $X \otimes Y$ (10)



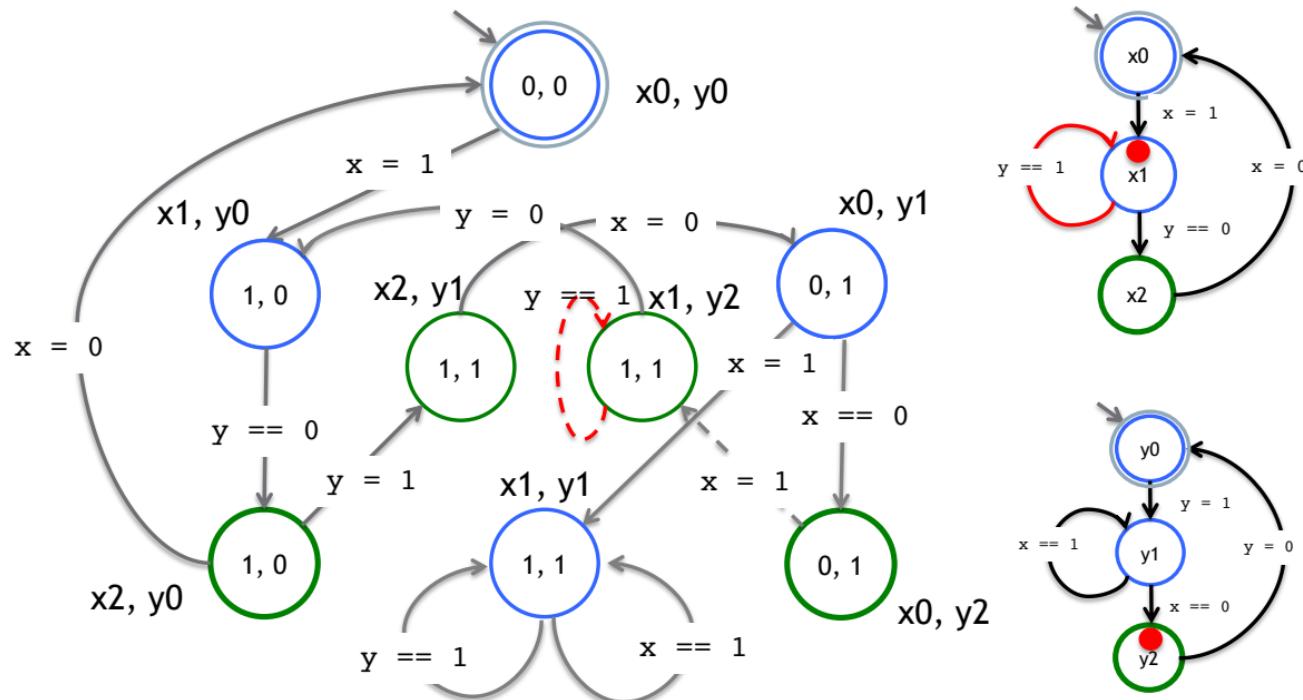
# Creating $X \otimes Y$ (11)



Hit a previously visited state: backtrack to  $x_1, y_2$



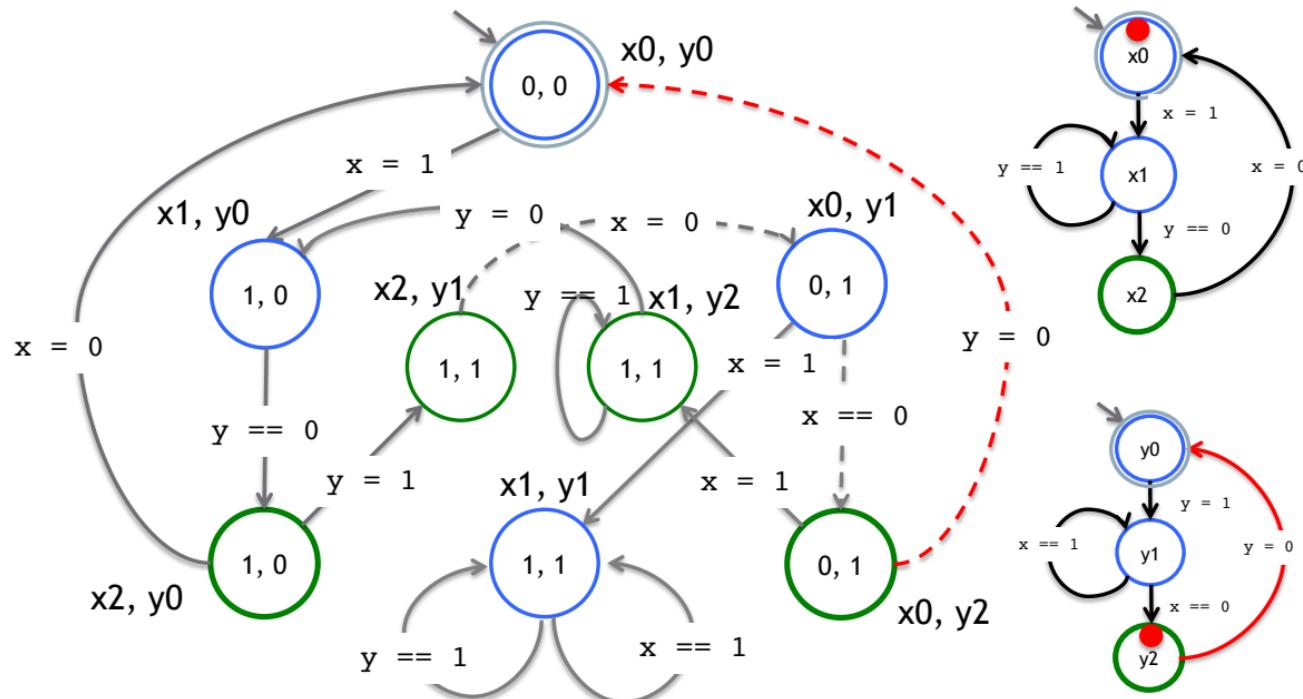
# Creating $X \otimes Y$ (12)



Hit a previously visited state  $x1, y2$ : backtrack to  $x1, y1$  and then to  $x0, y0$



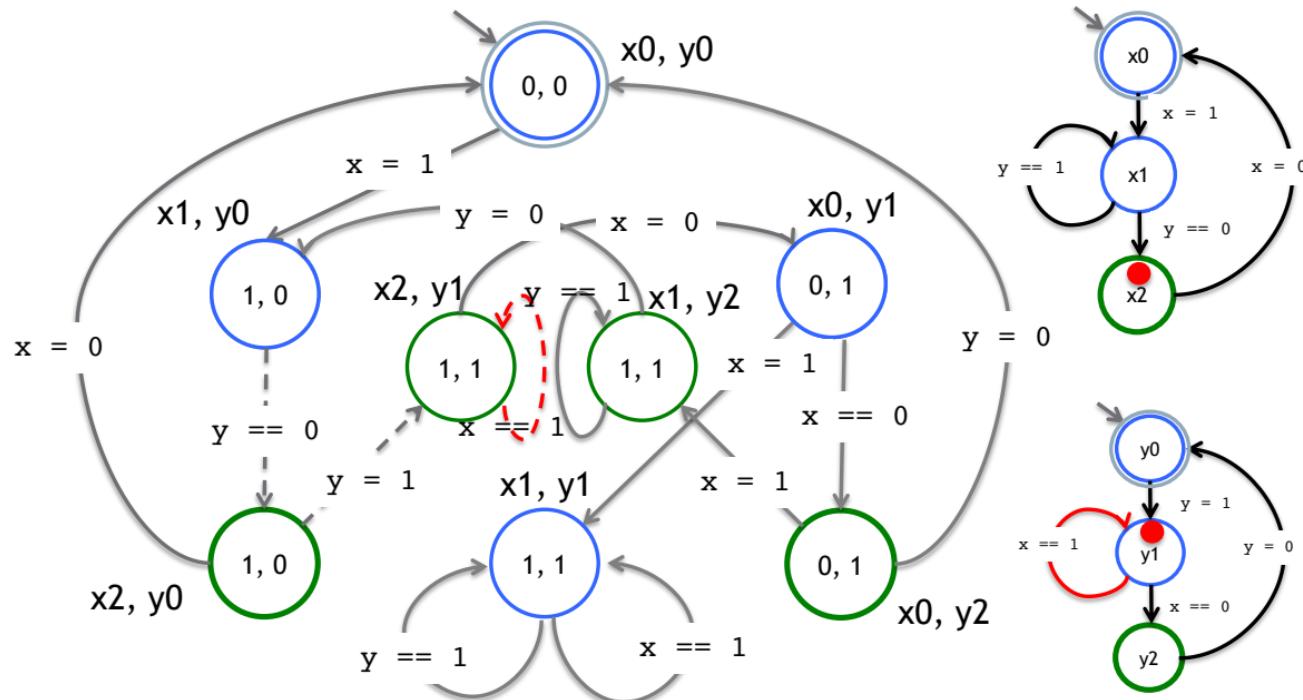
# Creating $X \otimes Y$ (13)



Hit a previously visited state  $x0, y0$ : backtrack to  $x0, y2$ ; then to  $x0, y1$ ; then  $x2, y1$



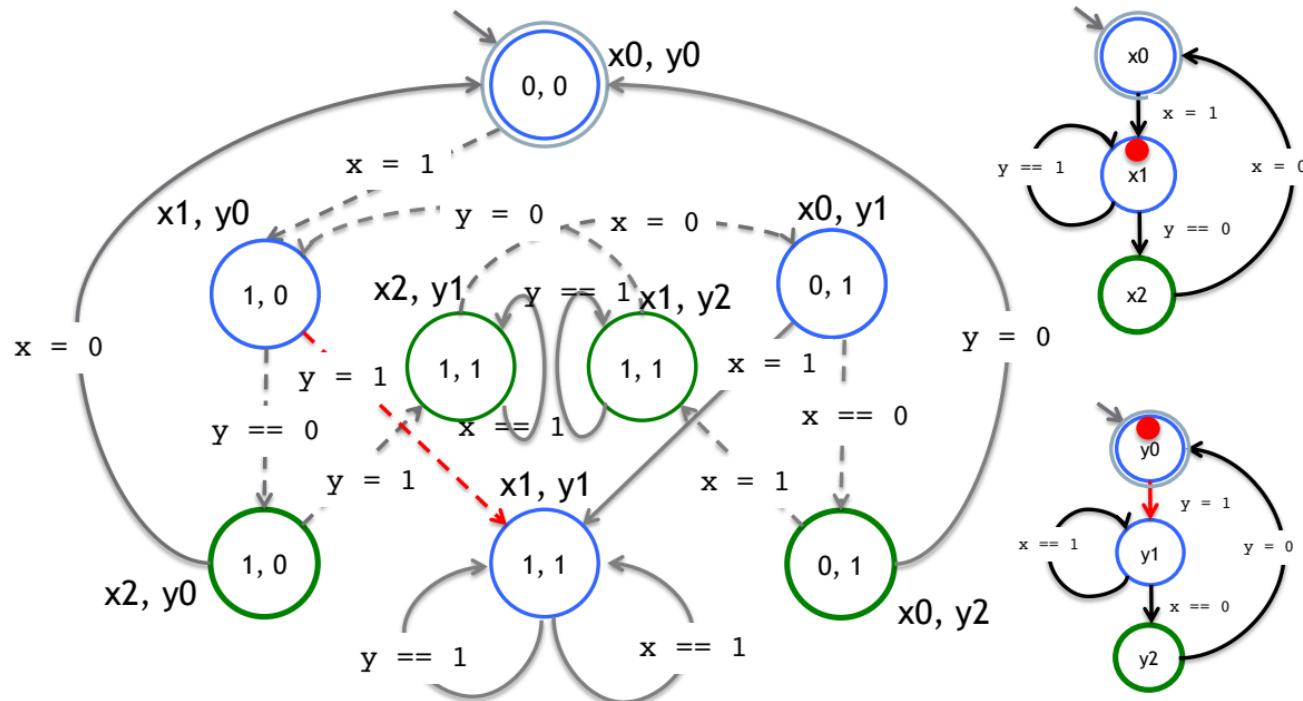
# Creating $X \otimes Y$ (14)



Hit a previously visited state  $x2, y1$ : backtrack to  $x2, y1$ ; then to  $x2, y0$ ; then  $x1, y0$



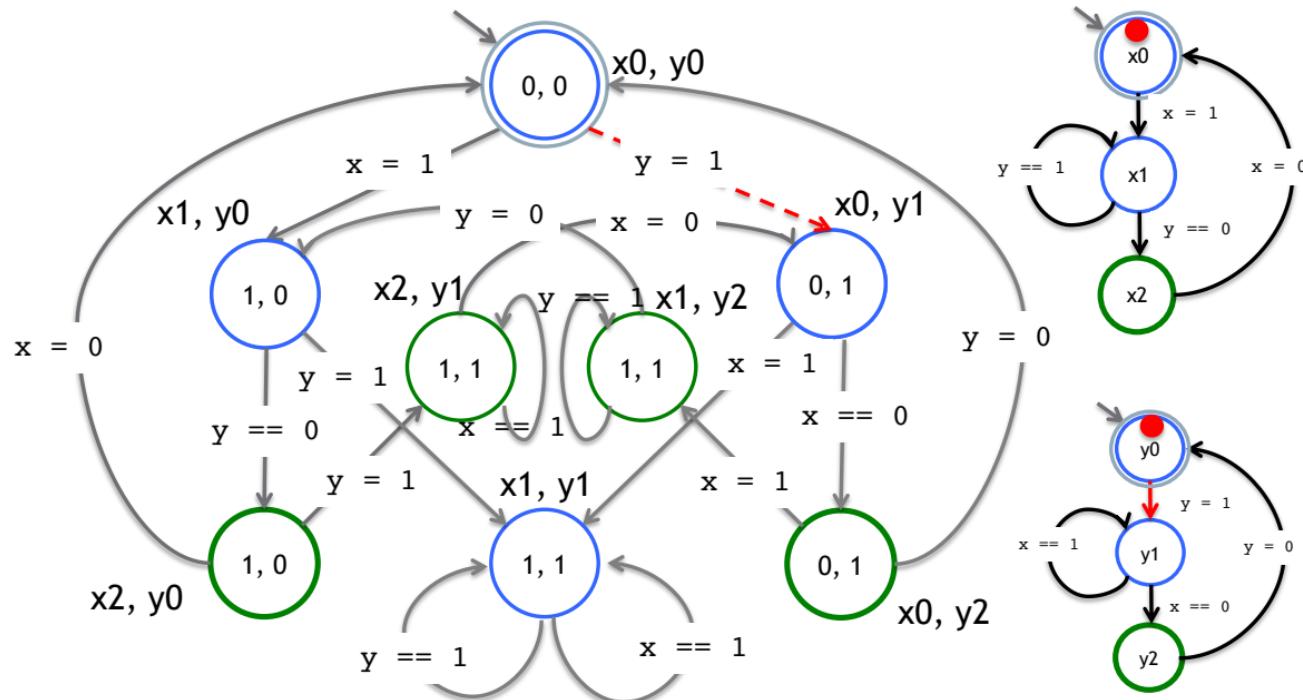
# Creating $X \otimes Y$ (15)



Hit a previously visited state  $x_1, y_1$ : backtrack to  $x_1, y_0$ ; then  $x_1, y_2$ ; then  $x_0, y_2$ ; .... all the way to  $x_0, y_0$ !



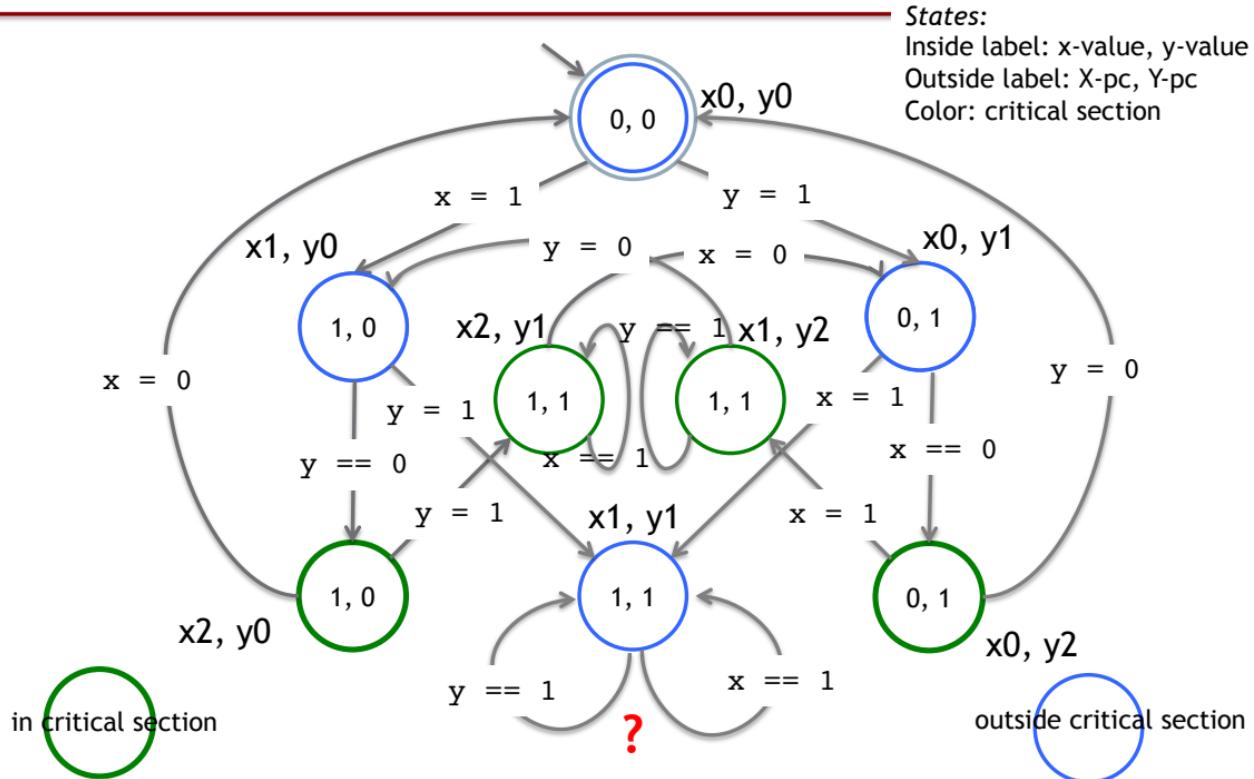
# Creating $X \otimes Y$ (16)



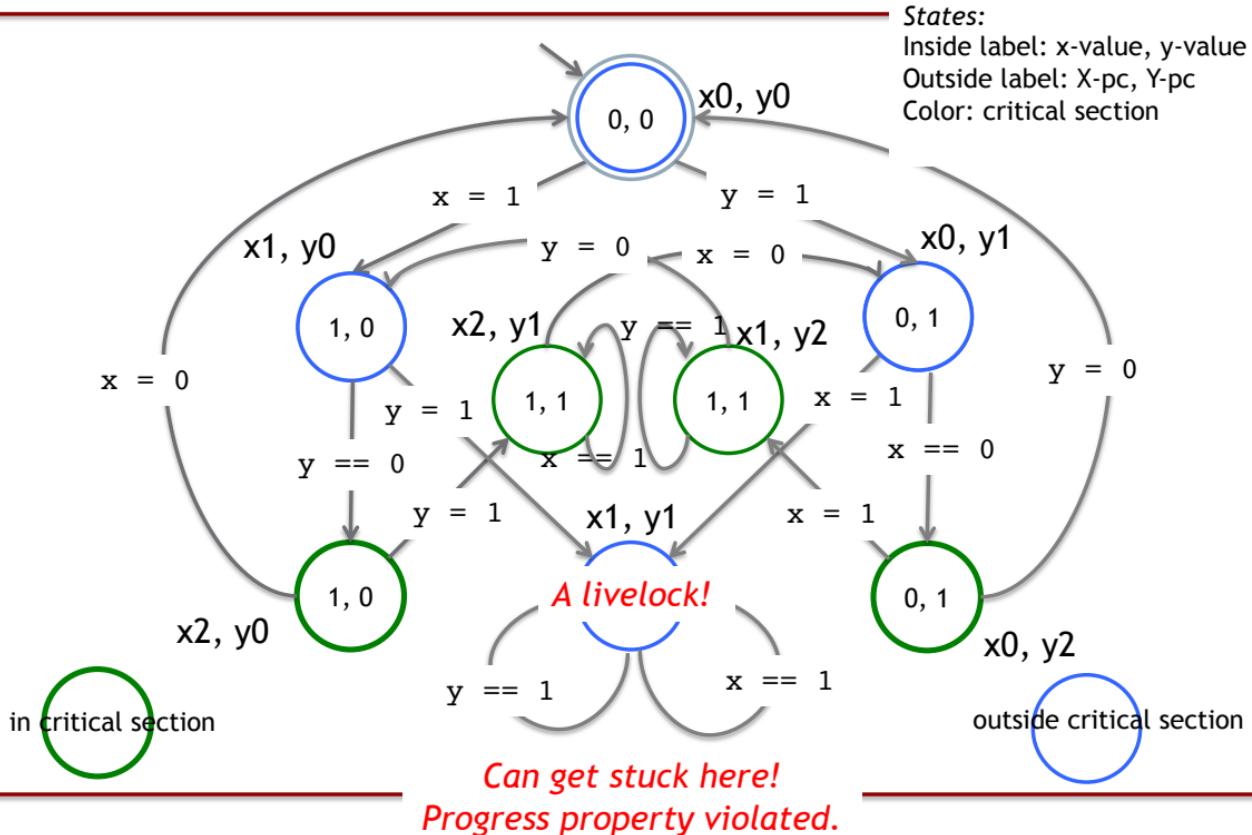
Hit a previously visited state  $x_0, y_1$ : backtrack to  $x_0, y_0$ ;  
No further transitions to explore: DONE!



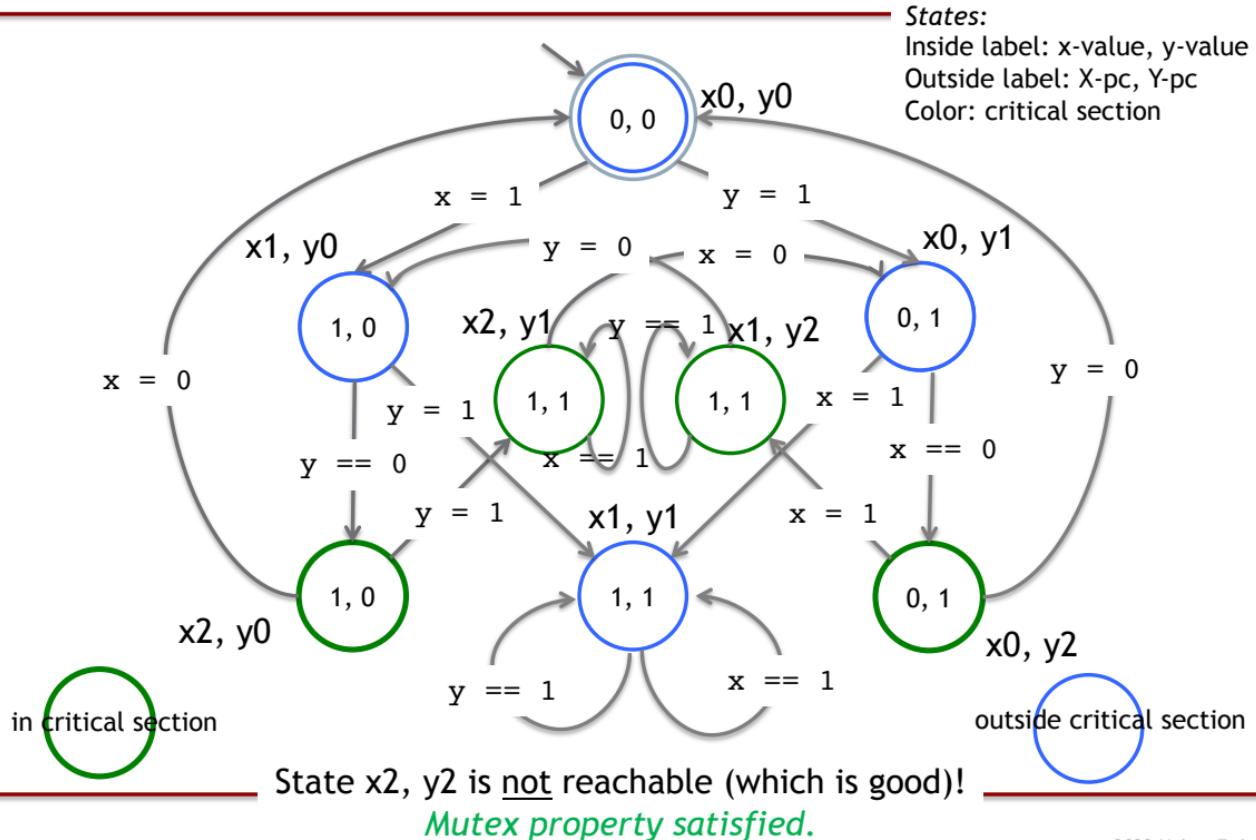
# The product FSA: final result



# The product FSA showing a livelock

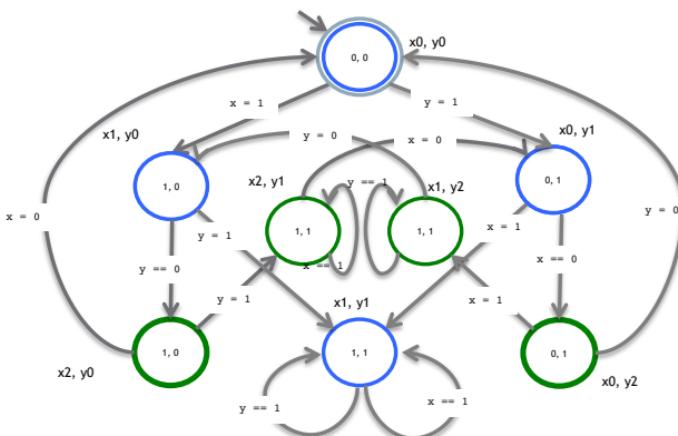


# The product FSA showing mutual exclusion



# The product FSA: properties summary

---



- Represents all possible global states: all possible, reachable interleavings of X and Y
  - Any path through this graph is a serialized execution path of the whole system
  - Includes only relevant detail for what we want to verify
  - **Progress** property: detects the livelock state (**fault**)
  - **Mutex** property: in every state, only one thread is in its critical section (state  $x_2, y_2$  is not reachable) (**good**)
-

# What did we just do?

---

## model checking!

```
stop = false;  
while (!stop) {  
    x = 1;  
    while (y == 1) {}  
    mutex++;  
    dSW(z);  
    mutex--;  
    x = 0;  
}
```

```
stop = false;  
while (!stop) {  
    y = 1;  
    while (x == 1) {}  
    mutex++;  
    dSW(z);  
    mutex--;  
    y = 0;  
}
```

**S** = Mutex-Program

**M** = X  $\otimes$  Y

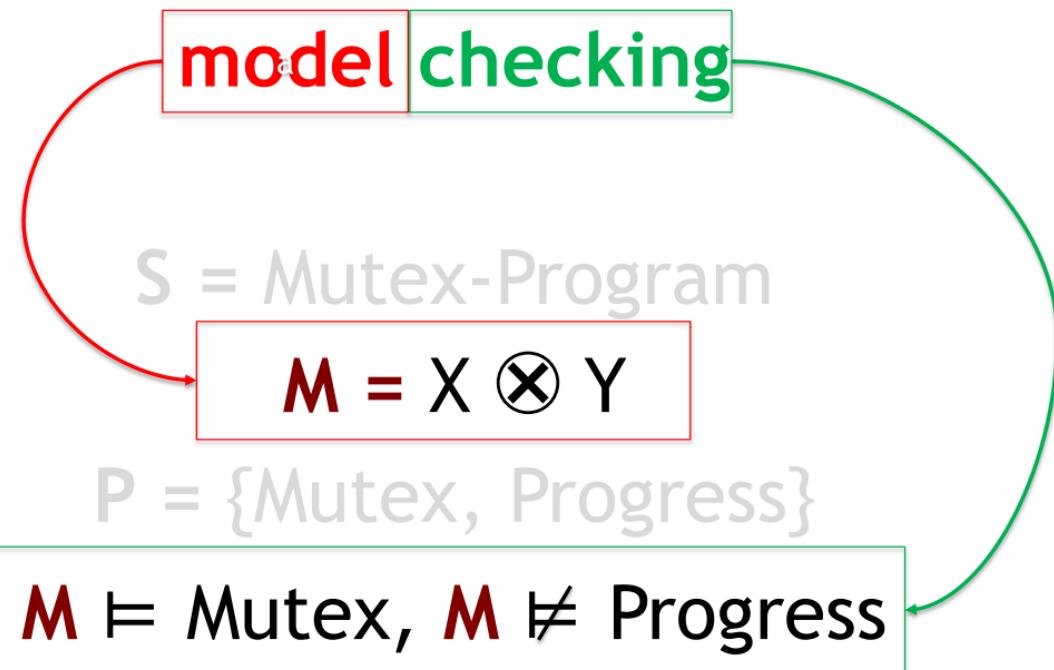
**P** = { Mutex, Progress }

**M**  $\models$  Mutex, **M**  $\not\models$  Progress, **M**  $\not\models$  P

---

# What did we just do?

---



---

[Clarke & Emerson 1981]:

“Model checking is an automated technique that, given a finite-state model of a system and a logical property, systematically checks whether this property holds for that model [starting from an initial state].”

Edmund Clarke, CMU & Allen Emerson, UT-Austin  
2007 Turing Award Winners

---

# Model checking is subject to limitations

---

- Models must be loyal/realistic enough
- Models should not restrict behavior with respect to properties
- Models must be finite-state
- Models must be sufficiently abstract (shouldn't have irrelevant details) so that model-checking is tractable



## We can apply similar process to much more complex models and properties

---

- Decide which details to eliminate/keep in system
  - Create the model in a formal modeling language with precise semantics
  - Express properties we want to verify using formal logic
  - Automatically generate the product FSA and check properties against this FSA
  - If product FSA is too big, manage the state space via further reasonable optimizations/approximations
-

## We can apply similar process to much more complex models and properties

- Decide which details to eliminate/keep in system
- Create the model in a formal modeling language with precise semantics (**Promela**)
- Express properties we want to verify using formal logic (**Linear Temporal Logic**)
- Automatically generate the product FSA and check properties against this FSA (**Spin**)
- If product FSA is too big, manage the state space via further reasonable optimizations/approximations (**various Spin options**)

# Model checking is not the only formal verification technique

---

- Abstract interpretation (already discussed)
- Symbolic execution
- Formal correctness proofs

# Formal verification techniques

---

- **Model checking** (what we started to describe)
    - Based on exploration of state space by a formal model of program
  - **Abstract interpretation:** execute program with abstract inputs (as in division-by-zero analysis) and explore its abstract state space
  - **Symbolic execution:** execute program with symbolic inputs (raw variables, no values) explore its symbolic state space (applied to automatic test case generation)
  - **Formal correctness proofs**
    - Based on formal program semantics: axiomatic, operational, denotational
    - *Tools:* logic(s), inference, induction
    - Often based on **Hoare Logic** via pre- and post-conditions
    - May operate on formal models in a specific notation or directly on source programs (Java, C)
-